

TRES Tidyverse Tutorial

Raphael, Pratik and Theo

2020-06-11

Contents

2	Outline	5
3	About	5
4	Schedule	5
5	Possible extras	6
6	Join	6
7	1 Reading files and string manipulation	7
8	1.1 Data import and export with <code>readr</code>	7
9	1.2 String manipulation with <code>stringr</code>	10
10	1.3 String interpolation with <code>glue</code>	19
11	1.4 Strings in <code>ggplot</code>	20
12	2 Reshaping data tables in the tidyverse, and other things	23
13	2.1 The new data frame: <code>tibble</code>	24
14	2.2 The concept of tidy data	26
15	2.3 Reshaping with <code>tidyr</code>	28
16	2.4 Extra: factors and the <code>forcats</code> package	36
17	2.5 External resources	42
18	3 Data manipulation with <code>dplyr</code>	43
19	3.1 Introduction	43
20	3.2 Working with existing variables	44
21	3.3 Working with observations	53
22	3.4 Making new variables	59
23	3.5 Working with multiple tables	68
24	4 Working with lists and iteration	73
25	4.1 List columns with <code>tidyr</code>	73
26	4.2 Iteration with <code>map</code>	76
27	4.3 More <code>map</code> variants	80
28	4.4 Other functions for working with lists	84
29	4.5 To add: <code>patchwork</code>	88

Outline

This is the readable version of the TRES tidyverse tutorial. A convenient PDF version can be downloaded by clicking the PDF document icon in the header bar.

About

The TRES tidyverse tutorial is an online workshop on how to use the tidyverse, a set of packages in the R computing language designed at making data handling and plotting easier.

This tutorial will take the form of a one hour per week video stream via Google Meet, every Friday morning at 10.00 (Groningen time) starting from the 29th of May, 2020 and lasting for a couple of weeks (depending on the number of topics we want to cover, but there should be at least 5).

PhD students from outside our department are welcome to attend.

Schedule

Topic	Package	Instructor	Date*
Reading data and string manipulation	readr, stringr, glue	Pratik	29/05/20
Data and reshaping	tibble, tidyr	Raphael	05/06/20
Manipulating data	dplyr	Theo	12/06/20
Working with lists and iteration	purrr	Pratik	19/06/20
Plotting	ggplot2	Raphael	26/06/20
Regular expressions	regex	Richel	03/07/20
Programming with the tidyverse	rlang	Pratik	10/07/20

44 Possible extras

- 45 • Reproducibility and package-making (with e.g. `usethis`)
- 46 • Embedding C++ code with `Rcpp`

48 Join

49 Join the Slack by clicking this link (Slack account required).

50 *Tentative dates.

Chapter 1

Reading files and string manipulation



Every use case is ridiculous
until it happens to you.

Load the packages for the day.

```
library(readr)
library(stringr)
library(glue)
```

1.1 Data import and export with readr

Data in the wild with which ecologists and evolutionary biologists deal is most often in the form of a text file, usually with the extensions `.csv` or `.txt`. Often,

such data has to be written to file from within R. `readr` contains a number of functions to help with reading and writing text files.

1.1.1 Reading data

Reading in a csv file with `readr` is done with the `read_csv` function, a faster alternative to the base R `read.csv`. Here, `read_csv` is applied to the `mtcars` example.

```
# get the filepath of the example
some_example = readr_example("mtcars.csv")

# read the file in
some_example = read_csv(some_example)

head(some_example)
#> # A tibble: 6 x 11
#>   mpg   cyl  disp    hp  drat    wt   qsec    vs  am  gear  carb
#>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1  21     6   160   110  3.9   2.62  16.5    0    1     4     4
#> 2  21     6   160   110  3.9   2.88  17.0    0    1     4     4
#> 3 22.8     4   108    93  3.85  2.32  18.6    1    1     4     1
#> 4 21.4     6   258   110  3.08  3.22  19.4    1    0     3     1
#> 5 18.7     8   360   175  3.15  3.44  17.0    0    0     3     2
#> 6 18.1     6   225   105  2.76  3.46  20.2    1    0     3     1
```

The `read_csv2` function is useful when dealing with files where the separator between columns is a semicolon `;`, and where the decimal point is represented by a comma `,`.

Other variants include:

- `read_tsv` for tab-separated files, and
- `read_delim`, a general case which allows the separator to be specified manually.

`readr` import function will attempt to guess the column type from the first N lines in the data. This N can be set using the function argument `guess_max`. The `n_max` argument sets the number of rows to read, while the `skip` argument sets the number of rows to be skipped before reading data.

By default, the column names are taken from the first row of the data, but they can be manually specified by passing a character vector to `col_names`.

There are some other arguments to the data import functions, but the defaults usually *just work*.

80 1.1.2 Writing data

81 Writing data uses the `write_*` family of functions, with implementations for
 82 `csv`, `csv2` etc. (represented by the asterisk), mirroring the import functions
 83 discussed above. `write_*` functions offer the `append` argument, which allow a
 84 data frame to be added to an existing file.

85 These functions are not covered here.

86 1.1.3 Reading and writing lines

87 Sometimes, there is text output generated in R which needs to be written to file,
 88 but is not in the form of a dataframe. A good example is model outputs. It is
 89 good practice to save model output as a text file, and add it to version control.
 90 Similarly, it may be necessary to import such text, either for display to screen,
 91 or to extract data.

92 This can be done using the `readr` functions `read_lines` and `write_lines`. Con-
 93 sider the model summary from a simple linear model.

```
# get the model
model = lm(mpg ~ wt, data = mtcars)
```

94 The model summary can be written to file. When writing lines to file, BE
 95 AWARE OF THE DIFFERENCES BETWEEN UNIX AND WINDOWS line
 96 separators. Usually, this causes no trouble.

```
# capture the model summary output
model_output = capture.output(summary(model))
```

```
# save it to file
write_lines(x = model_output,
  path = "model_output.txt")
```

97 This model output can be read back in for display, and each line of the model
 98 output is an element in a character vector.

```
# read in the model output and display
model_output = read_lines("model_output.txt")

# use cat to show the model output as it would be on screen
cat(model_output, sep = "\n")
#>
#> Call:
#> lm(formula = mpg ~ wt, data = mtcars)
#>
#> Residuals:
#>      Min       1Q   Median       3Q      Max
```

```

#> -4.543 -2.365 -0.125 1.410 6.873
#>
#> Coefficients:
#>             Estimate Std. Error t value Pr(>|t|)
#> (Intercept)  37.285      1.878   19.86 < 2e-16 ***
#> wt          -5.344      0.559   -9.56 1.3e-10 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 3.05 on 30 degrees of freedom
#> Multiple R-squared:  0.753, Adjusted R-squared:  0.745
#> F-statistic: 91.4 on 1 and 30 DF, p-value: 1.29e-10

```

99 These few functions demonstrate the most common uses of **readr**, but most
 100 other use cases for text data can be handled using different function arguments,
 101 including reading data off the web, unzipping compressed files before reading,
 102 and specifying the column types to control for type conversion errors.

103 Excel files

104 Finally, data is often shared or stored by well meaning people in the form
 105 of Microsoft Excel sheets. Indeed, Excel (especially when synced regularly to
 106 remote storage) is a good way of noting down observational data in the field.
 107 The **readxl** package allows importing from Excel files, including reading in
 108 specific sheets.

109 1.2 String manipulation with **stringr**

110 **stringr** is the tidyverse package for string manipulation, and exists in an in-
 111 teresting symbiosis with the **stringi** package. For the most part, **stringr** is a
 112 wrapper around **stringi**, and is almost always more than sufficient for day-to-day
 113 needs.

114 **stringr** functions begin with **str_**.

115 1.2.1 Putting strings together

116 Concatenate two strings with **str_c**, and duplicate strings with **str_dup**. Flat-
 117 ten a list or vector of strings using **str_flatten**.

```

# str_c works like paste(), choose a separator
str_c("this string", "this other string", sep = "_")
#> [1] "this string_this other string"

```

```

# str_dup works like rep
str_dup("this string", times = 3)
#> [1] "this stringthis stringthis string"

# str_flatten works on lists and vectors
str_flatten(string = as.list(letters), collapse = "_")
#> [1] "a_b_c_d_e_f_g_h_i_j_k_l_m_n_o_p_q_r_s_t_u_v_w_x_y_z"
str_flatten(string = letters, collapse = "-")
#> [1] "a-b-c-d-e-f-g-h-i-j-k-l-m-n-o-p-q-r-s-t-u-v-w-x-y-z"

118 str_flatten is especially useful when displaying the type of an object that
119 returns a list when class is called on it.

# get the class of a tibble and display it as a single string
class_tibble = class(tibble::tibble(a = 1))
str_flatten(string = class_tibble, collapse = ", ")
#> [1] "tbl_df, tbl, data.frame"

```

1.2.2 Detecting strings

Count the frequency of a pattern in a string with `str_count`. Returns an integer.
 Detect whether a pattern exists in a string with `str_detect`. Returns a logical
 and can be used as a predicate.

Both are vectorised, i.e, automatically applied to a vector of arguments.

```

# there should be 5 a-s here
str_count(string = "ababababa", pattern = "a")
#> [1] 5

# vectorise over the input string
# should return a vector of length 2, with integers 5 and 3
str_count(string = c("ababbababa", "banana"), pattern = "a")
#> [1] 5 3

# vectorise over the pattern to count both a-s and b-s
str_count(string = "ababababa", pattern = c("a", "b"))
#> [1] 5 4

```

Vectorising over both string and pattern works as expected.

```

# vectorise over both string and pattern
# counts a-s in first input, and b-s in the second
str_count(string = c("ababababa", "banana"),
          pattern = c("a", "b"))
#> [1] 5 1

# provide a longer pattern vector to search for both a-s

```

```

# and b-s in both inputs
str_count(string = c("ababababa", "banana"),
          pattern = c("a", "b",
                     "b", "a"))
#> [1] 5 1 4 3

```

126 **str_locate** locates the search pattern in a string, and returns the start and
127 end as a two column matrix.

```

# the behaviour of both str_locate and str_locate_all is
# to find the first match by default
str_locate(string = "banana", pattern = "ana")
#>      start end
#> [1,]      2  4

# str_detect detects a sequence in a string
str_detect(string = "Bananageddon is coming!",
           pattern = "na")
#> [1] TRUE

# str_detect is also vectorised and returns a two-element logical vector
str_detect(string = "Bananageddon is coming!",
           pattern = c("na", "don"))
#> [1] TRUE TRUE

# use any or all to convert a multi-element logical to a single logical
# here we ask if either of the patterns is detected
any(str_detect(string = "Bananageddon is coming!",
               pattern = c("na", "don")))
#> [1] TRUE

```

128 Detect whether a string starts or ends with a pattern. Also vectorised. Both
129 have a **negate** argument, which returns the negative, i.e., returns **FALSE** if the
130 search pattern is detected.

```

# taken straight from the examples, because they suffice
fruit <- c("apple", "banana", "pear", "pineapple")
# str_detect looks at the first character
str_starts(fruit, "p")
#> [1] FALSE FALSE TRUE TRUE

# str_ends looks at the last character
str_ends(fruit, "e")
#> [1] TRUE FALSE FALSE TRUE

# an example of negate = TRUE
str_ends(fruit, "e", negate = TRUE)
#> [1] FALSE TRUE TRUE FALSE

```

131 `str_subset` [WHICH IS NOT RELATED TO `str_sub`] helps with subsetting a
 132 character vector based on a `str_detect` predicate. In the example, all elements
 133 containing “banana” are subset.

134 `str_which` has the same logic except that it returns the vector position and not
 135 the elements.

```
# should return a subset vector containing the first two elements
str_subset(c("banana",
             "bananageddon is coming",
             "applegeddon is not real"),
           pattern = "banana")
#> [1] "banana"          "bananageddon is coming"

# returns an integer vector
str_which(c("banana",
            "bananageddon is coming",
            "applegeddon is not real"),
          pattern = "banana")
#> [1] 1 2
```

136 1.2.3 Matching strings

137 `str_match` returns all positive matches of the pattern in the string. The return
 138 type is a list, with one element per search pattern.

139 A simple case is shown below where the search pattern is the phrase “banana”.

```
str_match(string = c("banana",
                     "bananageddon",
                     "bananas are bad"),
          pattern = "banana")
#>      [,1]
#> [1,] "banana"
#> [2,] "banana"
#> [3,] "banana"
```

140 The search pattern can be extended to look for multiple subsets of the search
 141 pattern. Consider searching for dates and times.

142 Here, the search pattern is a `regex` pattern that looks for a set of four digits
 143 (`\\d{4}`) and a month name (`\\w+`) separated by a hyphen. There’s much more
 144 to be explored in dealing with dates and times in `lubridate`, another `tidyverse`
 145 package.

146 The return type is a list, each element is a character matrix where the first
 147 column is the string subset matching the full search pattern, and then as many
 148 columns as there are parts to the search pattern. The parts of interest in the

149 search pattern are indicated by wrapping them in parentheses. For example, in
 150 the case below, wrapping `[-.]` in parentheses will turn it into a distinct part
 151 of the search pattern.

```
# first with [-.] treated simply as a separator
str_match(string = c("1970-somemonth-01",
                     "1990-anothermonth-01",
                     "2010-thismonth-01"),
          pattern = "(\\d{4})[-.](\\w+)")
#>      [,1]      [,2] [,3]
#> [1,] "1970-somemonth" "1970" "somemonth"
#> [2,] "1990-anothermonth" "1990" "anothermonth"
#> [3,] "2010-thismonth" "2010" "thismonth"

# then with [-.] actively searched for
str_match(string = c("1970-somemonth-01",
                     "1990-anothermonth-01",
                     "2010-thismonth-01"),
          pattern = "(\\d{4})([-.])(\\w+)")
#>      [,1]      [,2] [,3] [,4]
#> [1,] "1970-somemonth" "1970" "-" "somemonth"
#> [2,] "1990-anothermonth" "1990" "-" "anothermonth"
#> [3,] "2010-thismonth" "2010" "-" "thismonth"
```

152 Multiple possible matches are dealt with using `str_match_all`. An example
 153 case is uncertainty in date-time in raw data, where the date has been entered
 154 as `1970-somemonth-01` or `1970/anothermonth/01`.

155 The return type is a list, with one element per input string. Each element is a
 156 character matrix, where each row is one possible match, and each column after
 157 the first (the full match) corresponds to the parts of the search pattern.

```
# first with a single date entry
str_match_all(string = c("1970-somemonth-01 or maybe 1990/anothermonth/01"),
              pattern = "(\\d{4})[\\-\\\\/](\\w+)")
#> [[1]]
#>      [,1]      [,2] [,3]
#> [1,] "1970-somemonth" "1970" "somemonth"
#> [2,] "1990/anothermonth" "1990" "anothermonth"

# then with multiple date entries
str_match_all(string = c("1970-somemonth-01 or maybe 1990/anothermonth/01",
                        "1990-somemonth-01 or maybe 2001/anothermonth/01"),
              pattern = "(\\d{4})[\\-\\\\/](\\w+)")
#> [[1]]
#>      [,1]      [,2] [,3]
#> [1,] "1970-somemonth" "1970" "somemonth"
#> [2,] "1990/anothermonth" "1990" "anothermonth"
```

```
#>
#> [[2]]
#>      [,1]      [,2]      [,3]
#> [1,] "1990-somemonth" "1990" "somemonth"
#> [2,] "2001/anothermonth" "2001" "anothermonth"
```

1.2.4 Simpler pattern extraction

The full functionality of `str_match_*` can be boiled down to the most common use case, extracting one or more full matches of the search pattern using `str_extract` and `str_extract_all` respectively.

`str_extract` returns a character vector with the same length as the input string vector, while `str_extract_all` returns a list, with a character vector whose elements are the matches.

```
# extracting the first full match using str_extract
str_extract(string = c("1970-somemonth-01 or maybe 1990/anothermonth/01",
                       "1990-somemonth-01 or maybe 2001/anothermonth/01"),
            pattern = "(\\d{4})[\\-\\/](\\[a-z\\]+)")
#> [1] "1970-somemonth" "1990-somemonth"

# extracting all full matches using str_extract_all
str_extract_all(string = c("1970-somemonth-01 or maybe 1990/anothermonth/01",
                           "1990-somemonth-01 or maybe 2001/anothermonth/01"),
                pattern = "(\\d{4})[\\-\\/](\\[a-z\\]+)")
#> [[1]]
#> [1] "1970-somemonth" "1990/anothermonth"
#>
#> [[2]]
#> [1] "1990-somemonth" "2001/anothermonth"
```

1.2.5 Breaking strings apart

`str_split`, `str_sub`, In the above date-time example, when reading filenames from a path, or when working sequences separated by a known pattern generally, `str_split` can help separate elements of interest.

The return type is a list similar to `str_match`.

```
# split on either a hyphen or a forward slash
str_split(string = c("1970-somemonth-01",
                     "1990/anothermonth/01"),
          pattern = "[\\-\\/]")
#> [[1]]
#> [1] "1970"      "somemonth" "01"
```

```
#>
#> [[2]]
#> [1] "1990"          "anothermonth" "01"
```

170 This can be useful in recovering simulation parameters from a filename, but may
 171 require some knowledge of `regex`.

```
# assume a simulation output file
filename = "sim_param1_0.01_param2_0.05_param3_0.01.ext"

# not quite there
str_split(filename, pattern = "_")
#> [[1]]
#> [1] "sim"      "param1"   "0.01"     "param2"   "0.05"     "param3"   "0.01.ext"

# not really
str_split(filename,
            pattern = "sim_")
#> [[1]]
#> [1] ""
#> [2] "param1_0.01_param2_0.05_param3_0.01.ext"

# getting there but still needs work
str_split(filename,
            pattern = "(sim_)|_*param\\d{1}|(.ext)")
#> [[1]]
#> [1] ""      ""      "0.01"  "0.05"  "0.01"  ""

172 str_split_fixed split the string into as many pieces as specified, and can be
173 especially useful dealing with filepaths.

# split on either a hyphen or a forward slash
str_split_fixed(string = "dir_level_1/dir_level_2/file.ext",
                pattern = "/",
                n = 2)
#>      [,1]      [,2]
#> [1,] "dir_level_1" "dir_level_2/file.ext"
```

174 1.2.6 Replacing string elements

175 `str_replace` is intended to replace the search pattern, and can be co-opted
 176 into the task of recovering simulation parameters or other data from regularly
 177 named files. `str_replace_all` works the same way but replaces all matches of
 178 the search pattern.

```
# replace all unwanted characters from this hypothetical filename with spaces
filename = "sim_param1_0.01_param2_0.05_param3_0.01.ext"
```



```

str_replace_all(filename,
                 pattern = "(sim_)|_*param\\d{1}_|(.ext)",
                 replacement = " ")
#> [1] " 0.01 0.05 0.01 "

```

179 `str_remove` is a wrapper around `str_replace` where the replacement is set to
 180 `" "`. This is not covered here.

181 Having replaced unwanted characters in the filename with spaces, `str_trim`
 182 offers a way to remove leading and trailing whitespaces.

```

# trim whitespaces from this filename after replacing unwanted text
filename = "sim_param1_0.01_param2_0.05_param3_0.01.ext"
filename_with_spaces = str_replace_all(filename,
                                       pattern = "(sim_)|_*param\\d{1}_|(.ext)",
                                       replacement = " ")

filename_without_spaces = str_trim(filename_with_spaces)
filename_without_spaces
#> [1] "0.01 0.05 0.01"

# the result can be split on whitespaces to return useful data
str_split(filename_without_spaces, " ")
#> [[1]]
#> [1] "0.01" "0.05" "0.01"

```

183 1.2.7 Subsetting within strings

184 When strings are highly regular, useful data can be extracted from a string using
 185 `str_sub`. In the date-time example, the year is always represented by the first
 186 four characters.

```

# get the year as characters 1 - 4
str_sub(string = c("1970-somemonth-01",
                  "1990-anothermonth-01",
                  "2010-thismonth-01"),
        start = 1, end = 4)
#> [1] "1970" "1990" "2010"

```

187 Similarly, it's possible to extract the last few characters using negative indices.

```

# get the day as characters -2 to -1
str_sub(string = c("1970-somemonth-01",
                  "1990-anothermonth-21",
                  "2010-thismonth-31"),
        start = -2, end = -1)
#> [1] "01" "21" "31"

```

188 Finally, it's also possible to replace characters within a string based on the
 189 position. This requires using the assignment operator `<-`.

```
# replace all days in these dates to 01
date_times = c("1970-somemonth-25",
               "1990-anothermonth-21",
               "2010-thismonth-31")

# a strictly necessary use of the assignment operator
str_sub(date_times,
        start = -2, end = -1) <- "01"

date_times
#> [1] "1970-somemonth-01" "1990-anothermonth-01" "2010-thismonth-01"
```

190 1.2.8 Padding and truncating strings

191 Strings included in filenames or plots are often of unequal lengths, especially
 192 when they represent numbers. `str_pad` can pad strings with suitable characters
 193 to maintain equal length filenames, with which it is easier to work.

```
# pad so all values have three digits
str_pad(string = c("1", "10", "100"),
        width = 3,
        side = "left",
        pad = "0")
#> [1] "001" "010" "100"
```

194 Strings can also be truncated if they are too long.

```
str_trunc(string = c("bananas are great and wonderful
                      and more stuff about bananas and
                      it really goes on about bananas"),
          width = 27,
          side = "right", ellipsis = "etc. etc.")
#> [1] "bananas are great etc. etc."
```

195 1.2.9 Stringr aspects not covered here

196 Some `stringr` functions are not covered here. These include:

- 197 • `str_wrap` (of dubious use),
- 198 • `str_interp`, `str_glue*` (better to use `glue`; see below),
- 199 • `str_sort`, `str_order` (used in sorting a character vector),
- 200 • `str_to_case*` (case conversion), and

- `str_view*` (a graphical view of search pattern matches).
 - `word`, `boundary` etc. The use of `word` is covered below.
- `stringi`, of which `stringr` is a wrapper, offers a lot more flexibility and control.

1.3 String interpolation with glue

The idea behind string interpolation is to procedurally generate new complex strings from pre-existing data.

`glue` is as simple as the example shown.

```
# print that each car name is a car model
cars = rownames(head(mtcars))
glue('The {cars} is a car model')
#> The Mazda RX4 is a car model
#> The Mazda RX4 Wag is a car model
#> The Datsun 710 is a car model
#> The Hornet 4 Drive is a car model
#> The Hornet Sportabout is a car model
#> The Valiant is a car model
```

This creates and prints a vector of car names stating each is a car model.

The related `glue_data` is even more useful in printing from a dataframe. In this example, it can quickly generate command line arguments or filenames.

```
# use dataframes for now
parameter_combinations = data.frame(param1 = letters[1:5],
                                     param2 = 1:5)

# for command line arguments or to start multiple job scripts on the cluster
glue_data(parameter_combinations,
           'simulation-name {param1} {param2}')
#> simulation-name a 1
#> simulation-name b 2
#> simulation-name c 3
#> simulation-name d 4
#> simulation-name e 5

# for filenames
glue_data(parameter_combinations,
           'sim_data_param1_{param1}_param2_{param2}.ext')
#> sim_data_param1_a_param2_1.ext
#> sim_data_param1_b_param2_2.ext
#> sim_data_param1_c_param2_3.ext
```

```
#> sim_data_param1_d_param2_4.ext
#> sim_data_param1_e_param2_5.ext
```

211 Finally, the convenient `glue_sql` and `glue_data_sql` are used to safely write
 212 SQL queries where variables from data are appropriately quoted. This is not
 213 covered here, but it is good to know it exists.

214 `glue` has some more functions — `glue_safe`, `glue_collapse`, and `glue_col`,
 215 but these are infrequently used. Their functionality can be found on the `glue`
 216 github page.

217 1.4 Strings in ggplot

218 `ggplot` has two `geoms` (wait for the `ggplot` tutorial to understand more about
 219 `geoms`) that work with text: `geom_text` and `geom_label`. These `geoms` allow
 220 text to be pasted on to the main body of a plot.

221 Often, these may overlap when the data are closely spaced. The pack-
 222 age `ggrepel` offers another `geom`, `geom_text_repel` (and the related
 223 `geom_label_repel`) that help arrange text on a plot so it doesn't over-
 224 lap with other features. This is *not perfect*, but it works more often than
 225 not.

226 More examples can be found on the `ggrepel` website.

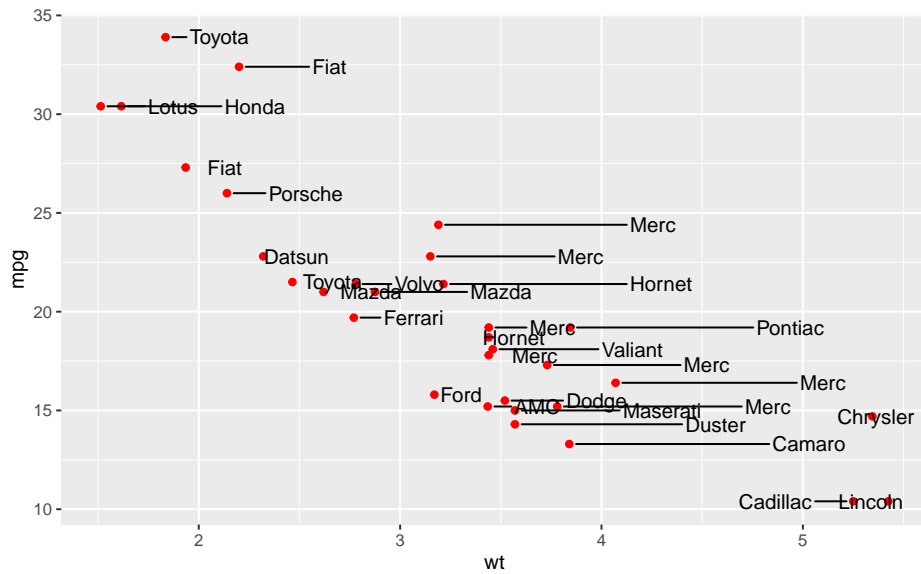
227 Here, the arguments to `geom_text_repel` are taken both from the `mtcars` data
 228 (position), as well as from the car brands extracted using the `stringr::word`
 229 (labels), which tries to separate strings based on a regular pattern.

230 The details of `ggplot` are covered in a later tutorial.

```
library(ggplot2)
library(ggrepel)

# prepare car labels using word function
car_labels = word(rownames(mtcars))

ggplot(mtcars,
  aes(x = wt, y = mpg,
    label = rownames(mtcars)))+
  geom_point(colour = "red")+
  geom_text_repel(aes(label = car_labels),
    direction = "x",
    nudge_x = 0.2,
    box.padding = 0.5,
    point.padding = 0.5)
```



231

232 This is not a good looking plot, because it breaks other rules of plot design,
 233 such as whether this sort of plot should be made at all. Labels and text need
 234 to be applied sparingly, for example drawing attention or adding information to
 235 outliers.

Chapter 2

Reshaping data tables in the tidyverse, and other things

Raphael Scherrer



Every use case is ridiculous
until it happens to you.

```
library(tibble)  
library(tidyr)
```

In this chapter we will learn what *tidy* means in the context of the tidyverse, and how to reshape our data into a tidy format using the `tidyr` package. But first, let us take a detour and introduce the `tibble`.

245

2.1 The new data frame: tibble

246 The **tibble** is the recommended class to use to store tabular data in the tidy-
 247 verse. Consider it as the operational unit of any data science pipeline. For most
 248 practical purposes, a **tibble** is basically a **data.frame**.

```
# Make a data frame
data.frame(who = c("Pratik", "Theo", "Raph"), chapt = c("1, 4", "3", "2, 5"))
#>      who chapt
#> 1 Pratik  1, 4
#> 2  Theo   3
#> 3  Raph  2, 5

# Or an equivalent tibble
tibble(who = c("Pratik", "Theo", "Raph"), chapt = c("1, 4", "3", "2, 5"))
#> # A tibble: 3 x 2
#>   who      chapt
#>   <chr>   <chr>
#> 1 Pratik 1, 4
#> 2 Theo   3
#> 3 Raph  2, 5
```

249 The difference between **tibble** and **data.frame** is in its display and in the way it
 250 is subsetting, among others. Most functions working with **data.frame** will work
 251 with **tibble** and vice versa. Use the **as*** family of functions to switch back and
 252 forth between the two if needed, using e.g. **as.data.frame** or **as_tibble**.

253 In terms of display, the tibble has the advantage of showing the class of each
 254 column: **chr** for **character**, **fct** for **factor**, **int** for **integer**, **dbl** for **numeric**
 255 and **lgl** for **logical**, just to name the main atomic classes. This may be more
 256 important than you think, because many hard-to-find bugs in R are due to
 257 wrong variable types and/or cryptic type conversions. This especially happens
 258 with **factor** and **character**, which can cause quite some confusion. More about
 259 this in the extra section at the end of this chapter!

260 Note that you can build a tibble by rows rather than by columns with **tribble**:

```
tribble(
  ~who, ~chapt,
  "Pratik", "1, 4",
  "Theo", "3",
  "Raph", "2, 5"
)
#> # A tibble: 3 x 2
#>   who      chapt
#>   <chr>   <chr>
#> 1 Pratik 1, 4
```



```
#> 2 Theo      3
#> 3 Raph     2, 5
```

As a rule of thumb, try to convert your tables to tibbles whenever you can, especially when the original table is *not* a data frame. For example, the principal component analysis function `prcomp` outputs a `matrix` of coordinates in principal component-space.

```
# Perform a PCA on mtcars
pca_scores <- prcomp(mtcars)$x
head(pca_scores) # looks like a data frame or a tibble...
#>
#>      PC1  PC2  PC3  PC4  PC5  PC6  PC7  PC8
#> Mazda RX4      -79.60  2.13 -2.15 -2.707 -0.702 -0.3149 -0.09870 -0.0779
#> Mazda RX4 Wag   -79.60  2.15 -2.22 -2.178 -0.884 -0.4534 -0.00355 -0.0957
#> Datsun 710      -133.89 -5.06 -2.14  0.346  1.106  1.1730  0.00576  0.1362
#> Hornet 4 Drive    8.52 44.99  1.23  0.827  0.424 -0.0579 -0.02431  0.2212
#> Hornet Sportabout 128.69 30.82  3.34 -0.521  0.737 -0.3329  0.10630 -0.0530
#> Valiant         -23.22 35.11 -3.26  1.401  0.803 -0.0884  0.23895  0.4239
#>
#>      PC9  PC10  PC11
#> Mazda RX4      -0.200 -0.2901  0.106
#> Mazda RX4 Wag   -0.353 -0.1928  0.107
#> Datsun 710      -0.198  0.0763  0.267
#> Hornet 4 Drive    0.356 -0.0906  0.209
#> Hornet Sportabout 0.153 -0.1886 -0.109
#> Valiant         0.101 -0.0377  0.276
class(pca_scores) # but is actually a matrix
#> [1] "matrix" "array"

# Convert to tibble
as_tibble(pca_scores)
#> # A tibble: 32 x 11
#>      PC1  PC2  PC3  PC4  PC5  PC6  PC7  PC8  PC9  PC10
#>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1  -79.6  2.13 -2.15 -2.71 -0.702 -0.315 -0.0987 -0.0779 -0.200 -0.290
#> 2  -79.6  2.15 -2.22 -2.18 -0.884 -0.453 -0.00355 -0.0957 -0.353 -0.193
#> 3 -134.  -5.06 -2.14  0.346  1.11  1.17  0.00576  0.136 -0.198  0.0763
#> 4   8.52 45.0  1.23  0.827  0.424 -0.0579 -0.0243  0.221  0.356 -0.0906
#> 5  129.  30.8  3.34 -0.521  0.737 -0.333  0.106  -0.0530  0.153 -0.189
#> 6  -23.2 35.1 -3.26  1.40  0.803 -0.0884  0.239  0.424  0.101 -0.0377
#> # ... with 26 more rows, and 1 more variable: PC11 <dbl>
```

This is important because a `matrix` can contain only one type of values (e.g. only `numeric` or `character`), while `tibble` (and `data.frame`) allow you to have columns of different types.

So, in the tidyverse we are going to work with tibbles, got it. But what does “tidy” mean exactly?

2.2 The concept of tidy data

When it comes to putting data into tables, there are many ways one could organize a dataset. The *tidy* format is one such format. According to the formal definition, a table is tidy if each column is a variable and each row is an observation. In practice, however, I found that this is not a very operational definition, especially in ecology and evolution where we often record multiple variables per individual. So, let's dig in with an example.

Say we have a dataset of several morphometrics measured on Darwin's finches in the Galapagos islands. Let's first get this dataset.

```
# We first simulate random data
beak_lengths <- rnorm(100, mean = 5, sd = 0.1)
beak_widths <- rnorm(100, mean = 2, sd = 0.1)
body_weights <- rgamma(100, shape = 10, rate = 1)
islands <- rep(c("Isabela", "Santa Cruz"), each = 50)

# Assemble into a tibble
data <- tibble(
  id = 1:100,
  body_weight = body_weights,
  beak_length = beak_lengths,
  beak_width = beak_widths,
  island = islands
)

# Snapshot
data
#> # A tibble: 100 x 5
#>   id body_weight beak_length beak_width island
#>   <int>      <dbl>      <dbl>      <dbl> <chr>
#> 1     1      10.8        4.94      1.94 Isabela
#> 2     2      15.4        5.02      2.00 Isabela
#> 3     3      15.0        4.92      1.91 Isabela
#> 4     4       8.51        5.16      2.02 Isabela
#> 5     5      14.9        5.03      1.93 Isabela
#> 6     6       8.41        4.92      2.18 Isabela
#> # ... with 94 more rows
```

Here, we pretend to have measured `beak_length`, `beak_width` and `body_weight` on 100 birds, 50 of them from Isabela and 50 of them from Santa Cruz. In this tibble, each row is an individual bird. This is probably the way most scientists would record their data in the field. However, a single bird is not an “observation” in the sense used in the tidyverse. Our dataset is not tidy but *messy*.

The tidy equivalent of this dataset would be:

```

data <- pivot_longer(
  data,
  cols = c("body_weight", "beak_length", "beak_width"),
  names_to = "variable"
)
data
#> # A tibble: 300 x 4
#>   id island variable    value
#>   <int> <chr>   <chr>    <dbl>
#> 1     1  Isabela body_weight 10.8
#> 2     1  Isabela beak_length  4.94
#> 3     1  Isabela beak_width  1.94
#> 4     2  Isabela body_weight 15.4
#> 5     2  Isabela beak_length  5.02
#> 6     2  Isabela beak_width  2.00
#> # ... with 294 more rows

```

285 where each *measurement* (and not each *individual*) is now the unit of observation
 286 (the rows). The `pivot_longer` function is the easiest way to get to this format.
 287 It belongs to the `tidyr` package, which we'll cover in a minute.

288 As you can see our tibble now has three times as many rows and fewer columns.
 289 This format is rather unintuitive and not optimal for display. However, it pro-
 290 vides a very standardized and consistent way of organizing data that will be
 291 understood (and expected) by pretty much all functions in the tidyverse. This
 292 makes the tidyverse tools work well together and reduces the time you would
 293 otherwise spend reformatting your data from one tool to the next.

294 That does not mean that the *messy* format is useless though. There may be
 295 use-cases where you need to switch back and forth between formats. For this
 296 reason I prefer referring to these formats using their other names: *long* (tidy)
 297 versus *wide* (messy). For example, matrix operations work much faster on wide
 298 data, and the wide format arguably looks nicer for display. Luckily the `tidyr`
 299 package gives us the tools to reshape our data as needed, as we shall see shortly.

300 Another common example of wide-or-long dilemma is when dealing with *con-*
 301 *tingency tables*. This would be our case, for example, if we asked how many
 302 observations we have for each morphometric and each island. We use `table`
 303 (from base R) to get the answer:

```

# Make a contingency table
ctg <- with(data, table(island, variable))
ctg
#>           variable
#> island  beak_length beak_width body_weight
#> Isabela           50          50          50
#> Santa Cruz        50          50          50

```

304 A variety of statistical tests can be used on contingency tables such as Fisher's

exact test, the chi-square test or the binomial test. Contingency tables are in the wide format by construction, but they too can be pivoted to the long format, and the tidyverse manipulation tools will expect you to do so. Actually, `tibble` knows that very well and does it by default if you convert your `table` into a `tibble`:

```
# Contingency table is pivoted to the long-format automatically
as_tibble(ctg)
#> # A tibble: 6 x 3
#>   island      variable      n
#>   <chr>      <chr>    <int>
#> 1 Isabela  beak_length    50
#> 2 Santa Cruz beak_length    50
#> 3 Isabela  beak_width    50
#> 4 Santa Cruz beak_width    50
#> 5 Isabela  body_weight    50
#> 6 Santa Cruz body_weight    50
```

Summary: Tidy or not tidy

To sum up, the definition of what is tidy and what is not is somewhat subjective. Tables can be in long or wide format, and depending on the complexity of a dataset, there may even be some intermediate states. To be clear, the tidyverse does not only accept long tables, and wide tables may sometimes be the way to go. This is very use-case specific. Have a clear idea of what you want to do with your data (what tidyverse tools you will use), and use that to figure which format makes more sense. And remember, `tidyr` is here to easily do the switching for you.

2.3 Reshaping with `tidyr`

The `tidyr` package implements tools to easily switch between layouts and also perform a few other reshaping operations. Old school R users will be familiar with the `reshape` and `reshape2` packages, of which `tidyr` is the tidyverse equivalent. Beware that `tidyr` is about playing with the general *layout* of the dataset, while *operations* and *transformations* of the data are within the scope of the `dplyr` and `purrr` packages. All these packages work hand-in-hand really well, and analysis pipelines usually involve all of them. But today, we focus on the first member of this holy trinity, which is often the first one you'll need because you will want to reshape your data before doing other things. So, please hold your non-layout-related questions for the next chapters.

2.3.1 Pivoting

Pivoting a dataset between the long and wide layout is the main purpose of **tidyr** (check out the package's logo). We already saw the **pivot_longer** function above. This function converts a table from wide to long format. Similarly, there is a **pivot_wider** function that does exactly the opposite and takes you back to the wide format:

```
pivot_wider(
  data,
  names_from = "variable",
  values_from = "value",
  id_cols = c("id", "island")
)
#> # A tibble: 100 x 5
#>       id island  body_weight beak_length beak_width
#>   <int> <chr>      <dbl>      <dbl>      <dbl>
#> 1     1  Isabela    10.8        4.94        1.94
#> 2     2  Isabela    15.4        5.02        2.00
#> 3     3  Isabela    15.0        4.92        1.91
#> 4     4  Isabela     8.51        5.16        2.02
#> 5     5  Isabela    14.9        5.03        1.93
#> 6     6  Isabela     8.41        4.92        2.18
#> # ... with 94 more rows
```

The order of the columns is not exactly as it was, but this should not matter in a data analysis pipeline where you should access columns by their names. It is straightforward to change the order of the columns, but this is more within the scope of the **dplyr** package.

If you are familiar with earlier versions of the tidyverse, **pivot_longer** and **pivot_wider** are the respective equivalents of **gather** and **spread**, which are now deprecated.

There are a few other reshaping operations from **tidyr** that are worth knowing.

2.3.2 Handling missing values

Say we have some missing measurements in the column “value” of our finch dataset:

```
# We replace 100 random observations by NAs
ii <- sample(nrow(data), 100)
data$value[ii] <- NA
data
#> # A tibble: 300 x 4
#>       id island variable  value
```

```

#>   <int> <chr>   <chr>       <dbl>
#> 1     1  Isabela body_weight 10.8
#> 2     1  Isabela beak_length NA
#> 3     1  Isabela beak_width  NA
#> 4     2  Isabela body_weight NA
#> 5     2  Isabela beak_length 5.02
#> 6     2  Isabela beak_width  NA
#> # ... with 294 more rows

```

338 We could get rid of the rows that have missing values using `drop_na`:

```

drop_na(data, value)
#> # A tibble: 200 x 4
#>   id island variable    value
#>   <int> <chr>   <chr>      <dbl>
#> 1     1  Isabela body_weight 10.8
#> 2     2  Isabela beak_length 5.02
#> 3     3  Isabela body_weight 15.0
#> 4     3  Isabela beak_length 4.92
#> 5     4  Isabela body_weight 8.51
#> 6     4  Isabela beak_width 2.02
#> # ... with 194 more rows

```

339 Else, we could replace the NAs with some user-defined value:

```

replace_na(data, replace = list(value = -999))
#> # A tibble: 300 x 4
#>   id island variable    value
#>   <int> <chr>   <chr>      <dbl>
#> 1     1  Isabela body_weight 10.8
#> 2     1  Isabela beak_length -999
#> 3     1  Isabela beak_width  -999
#> 4     2  Isabela body_weight -999
#> 5     2  Isabela beak_length 5.02
#> 6     2  Isabela beak_width  -999
#> # ... with 294 more rows

```

340 where the `replace` argument takes a named list, and the names should refer to
 341 the columns to apply the replacement to.

342 We could also replace NAs with the most recent non-NA values:

```

fill(data, value)
#> # A tibble: 300 x 4
#>   id island variable    value
#>   <int> <chr>   <chr>      <dbl>
#> 1     1  Isabela body_weight 10.8
#> 2     1  Isabela beak_length 10.8
#> 3     1  Isabela beak_width 10.8

```

```
#> 4      2 Isabela body_weight 10.8
#> 5      2 Isabela beak_length  5.02
#> 6      2 Isabela beak_width   5.02
#> # ... with 294 more rows
```

Note that most functions in the tidyverse take a tibble as their first argument, and columns to which to apply the functions are usually passed as “objects” rather than character strings. In the above example, we passed the `value` column as `value`, not `"value"`. These column-objects are called by the tidyverse functions *in the context* of the data (the tibble) they belong to.

2.3.3 Splitting and combining cells

The `tidyr` package offers tools to split and combine columns. This is a nice extension to the string manipulations we saw last week in the `stringr` tutorial.

Say we want to add the specific dates when we took measurements on our birds (we would normally do this using `dplyr` but for now we will stick to the old way):

```
# Sample random dates for each observation
data$day <- sample(30, nrow(data), replace = TRUE)
data$month <- sample(12, nrow(data), replace = TRUE)
data$year <- sample(2019:2020, nrow(data), replace = TRUE)
data
#> # A tibble: 300 x 7
#>       id island variable    value  day month  year
#>   <int> <chr>   <chr>    <dbl> <int> <int> <int>
#> 1     1  Isabela body_weight 10.8     8     7  2020
#> 2     1  Isabela beak_length NA      19     7  2019
#> 3     1  Isabela beak_width  NA     17    12  2019
#> 4     2  Isabela body_weight NA      20    12  2020
#> 5     2  Isabela beak_length 5.02    21    10  2020
#> 6     2  Isabela beak_width  NA     23     2  2020
#> # ... with 294 more rows
```

We could combine the `day`, `month` and `year` columns into a single `date` column, with a dash as a separator, using `unite`:

```
data <- unite(data, day, month, year, col = "date", sep = "-")
data
#> # A tibble: 300 x 5
#>       id island variable    value date
#>   <int> <chr>   <chr>    <dbl> <chr>
#> 1     1  Isabela body_weight 10.8 8-7-2020
#> 2     1  Isabela beak_length NA   19-7-2019
#> 3     1  Isabela beak_width  NA   17-12-2019
```

```
#> 4      2 Isabela body_weight NA      20-12-2020
#> 5      2 Isabela beak_length 5.02 21-10-2020
#> 6      2 Isabela beak_width NA      23-2-2020
#> # ... with 294 more rows
```

356 Of course, we can revert back to the previous dataset by splitting the `date`
 357 column with `separate`.

```
separate(data, date, into = c("day", "month", "year"))
#> # A tibble: 300 x 7
#>       id island variable    value day month year
#>   <int> <chr>   <chr>      <dbl> <chr> <chr> <chr>
#> 1     1 Isabela body_weight 10.8   8     7    2020
#> 2     1 Isabela beak_length NA     19    7    2019
#> 3     1 Isabela beak_width NA     17   12    2019
#> 4     2 Isabela body_weight NA     20   12    2020
#> 5     2 Isabela beak_length 5.02  21   10    2020
#> 6     2 Isabela beak_width NA     23    2    2020
#> # ... with 294 more rows
```

358 But note that the `day`, `month` and `year` columns are now of class `character` and
 359 not `integer` anymore. This is because they result from the splitting of `date`,
 360 which itself was a `character` column.

361 You can also separate a single column into multiple *rows* using `separate_rows`:

```
separate_rows(data, date)
#> # A tibble: 900 x 5
#>       id island variable    value date
#>   <int> <chr>   <chr>      <dbl> <chr>
#> 1     1 Isabela body_weight 10.8   8
#> 2     1 Isabela body_weight 10.8   7
#> 3     1 Isabela body_weight 10.8 2020
#> 4     1 Isabela beak_length NA     19
#> 5     1 Isabela beak_length NA     7
#> 6     1 Isabela beak_length NA    2019
#> # ... with 894 more rows
```

362 2.3.4 Expanding tables using combinations

363 Instead of getting rid of rows with NAs, we may want to add rows with NAs,
 364 for example, for combinations of parameters that we did not measure.

```
data <- separate(data, date, into = c("day", "month", "year"))
to_rm <- with(data, island == "Santa Cruz" & year == "2020")
data <- data[!to_rm,]
tail(data)
#> # A tibble: 6 x 7
```



```
#>      id island      variable      value day  month year
#>   <int> <chr>      <chr>      <dbl> <chr> <chr> <chr>
#> 1    98 Santa Cruz beak_length  4.94 22    12    2019
#> 2    98 Santa Cruz beak_width   1.90 9     1     2019
#> 3    99 Santa Cruz body_weight 15.0 16     7     2019
#> 4    99 Santa Cruz beak_length NA    26    10    2019
#> 5    99 Santa Cruz beak_width   2.04 30     7     2019
#> 6   100 Santa Cruz beak_width   NA    23     3     2019
```

365 We could generate a tibble with all combinations of island, morphometric and
 366 year using `expand_grid`:

```
expand_grid(
  island = c("Isabela", "Santa Cruz"),
  year = c("2019", "2020")
)
#> # A tibble: 4 x 2
#>   island      year
#>   <chr>      <chr>
#> 1 Isabela    2019
#> 2 Isabela    2020
#> 3 Santa Cruz 2019
#> 4 Santa Cruz 2020
```

367 If we already have a tibble to work from that contains the variables to combine,
 368 we can use `expand` on that tibble:

```
expand(data, island, year)
#> # A tibble: 4 x 2
#>   island      year
#>   <chr>      <chr>
#> 1 Isabela    2019
#> 2 Isabela    2020
#> 3 Santa Cruz 2019
#> 4 Santa Cruz 2020
```

369 As you can see, we get all the combinations of the variables of interest, even
 370 those that are missing. But sometimes you might be interested in variables
 371 that are *nested* within each other and not *crossed*. For example, say we have
 372 measured birds at different locations within each island:

```
nrow_Isabela <- with(data, length(which(island == "Isabela")))
nrow_SantaCruz <- with(data, length(which(island == "Santa Cruz")))
sites_Isabela <- sample(c("A", "B"), size = nrow_Isabela, replace = TRUE)
sites_SantaCruz <- sample(c("C", "D"), size = nrow_SantaCruz, replace = TRUE)
sites <- c(sites_Isabela, sites_SantaCruz)
data$site <- sites
data
#> # A tibble: 232 x 8
```

```
#>      id island variable    value day  month year  site
#>   <int> <chr>   <chr>      <dbl> <chr> <chr> <chr> <chr>
#> 1      1  Isabela body_weight 10.8   8      7    2020  A
#> 2      1  Isabela beak_length NA     19     7    2019  B
#> 3      1  Isabela beak_width  NA     17     12   2019  B
#> 4      2  Isabela body_weight NA     20     12   2020  A
#> 5      2  Isabela beak_length 5.02  21     10   2020  A
#> 6      2  Isabela beak_width  NA     23     2    2020  A
#> # ... with 226 more rows
```

Of course, if sites A and B are on Isabela, they cannot be on Santa Cruz, where we have sites C and D instead. It would not make sense to **expand** assuming that **island** and **site** are crossed, instead, they are nested. We can therefore expand using the **nesting** function:

```
expand(data, nesting(island, site, year))
#> # A tibble: 6 x 3
#>   island    site year
#>   <chr>    <chr> <chr>
#> 1 Isabela    A    2019
#> 2 Isabela    A    2020
#> 3 Isabela    B    2019
#> 4 Isabela    B    2020
#> 5 Santa Cruz C    2019
#> 6 Santa Cruz D    2019
```

But now the missing data for Santa Cruz in 2020 are not accounted for because **expand** thinks the **year** is also nested within **island**. To get back the missing combination, we use **crossing**, the complement of **nesting**:

```
expand(data, crossing(nesting(island, site), year)) # both can be used together
#> # A tibble: 8 x 3
#>   island    site year
#>   <chr>    <chr> <chr>
#> 1 Isabela    A    2019
#> 2 Isabela    A    2020
#> 3 Isabela    B    2019
#> 4 Isabela    B    2020
#> 5 Santa Cruz C    2019
#> 6 Santa Cruz C    2020
#> # ... with 2 more rows
```

Here, we specify that **site** is nested within **island** and these two are crossed with **year**. Easy!

But wait a minute. These combinations are all very good, but our measurements have disappeared! We can get them back by levelling up to the **complete** function instead of using **expand**:

```

tail(complete(data, crossing(nesting(island, site), year)))
#> # A tibble: 6 x 8
#>   island      site year   id variable    value day  month
#>   <chr>      <chr> <chr> <int> <chr>      <dbl> <chr> <chr>
#> 1 Santa Cruz D    2019   95 beak_width  NA    13    10
#> 2 Santa Cruz D    2019   98 beak_length 4.94  22    12
#> 3 Santa Cruz D    2019   99 body_weight 15.0  16     7
#> 4 Santa Cruz D    2019   99 beak_length NA    26    10
#> 5 Santa Cruz D    2019   99 beak_width  2.04  30     7
#> 6 Santa Cruz D    2020   NA <NA>      NA    <NA> <NA>
# the last row has been added, full of NAs

```

385 which nicely keeps the rest of the columns in the tibble and just adds the missing
 386 combinations.

387 2.3.5 Nesting

388 The `tidyr` package has yet another feature that makes the tidyverse very pow-
 389 erful: the `nest` function. However, it makes little sense without combining it
 390 with the functions in the `purrr` package, so we will not cover it in this chapter
 391 but rather in the `purrr` chapter.

392 2.3.6 What else can be tidied up?

393 2.3.6.1 Model output with broom

394 Check out the `broom` package and its `tidy` function to tidy up messy linear
 395 model output, e.g.

```

library(broom)
fit <- lm(mpg ~ cyl, mtcars)
summary(fit)
#>
#> Call:
#> lm(formula = mpg ~ cyl, data = mtcars)
#>
#> Residuals:
#>    Min       1Q   Median       3Q      Max
#> -4.981 -2.119  0.222  1.072  7.519
#>
#> Coefficients:
#>              Estimate Std. Error t value Pr(>|t|)
#> (Intercept)   37.885      2.074   18.27 < 2e-16 ***
#> cyl           -2.876      0.322   -8.92 6.1e-10 ***
#> ---

```

```

#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 3.21 on 30 degrees of freedom
#> Multiple R-squared:  0.726, Adjusted R-squared:  0.717
#> F-statistic: 79.6 on 1 and 30 DF,  p-value: 6.11e-10
tidy(fit) # returns a tibble
#> # A tibble: 2 x 5
#>   term          estimate std.error statistic  p.value
#>   <chr>          <dbl>     <dbl>     <dbl>    <dbl>
#> 1 (Intercept)    37.9       2.07      18.3 8.37e-18
#> 2 cyl          -2.88       0.322     -8.92 6.11e-10

```

396 The **broom** package is just one package among a series of packages together
 397 known as **tidymodels** that deal with statistical models according to the tidyverse
 398 philosophy, and those include machine learning models.

399 2.3.6.2 Graphs with **tidygraph**

400 For some datasets, sometimes there is no trivial and intuitive way to store them
 401 into a table. This is the case, for example, for data underlying graphs (as in
 402 networks), which contain information about relations between entities. What
 403 is the unit of observation in a network? A node? An edge between two nodes?
 404 Nodes and edges in a network may each have node- or edge-specific variables
 405 mapped to them, and both may be equally valid units of observation. The
 406 **tidygraph** package has tools to store graph-data in a tidyverse-friendly object,
 407 consisting of two tibbles: one for node-specific information, the other for edge-
 408 specific information. This package goes hand in hand with the **ggraph**, that
 409 makes plotting networks compatible with the grammar of graphics.

410 2.3.6.3 Trees with **tidytree**

411 Phylogenetic trees are a special type of graphs suffering from the same issue,
 412 i.e. of being non-trivial to store in a table. The **tidytree** package and its
 413 companion **treeio** offer an interface to convert tree-like objects (from most
 414 format used by other packages and software) into a tidyverse-friendly format.
 415 Again, the point is that the rest of the tidyverse can be used to wrangle or plot
 416 this type of data in the same way as one would do with regular tabular data.
 417 For plotting a **tidytree** with the grammar of graphics, see **ggtree**.

418 2.4 Extra: factors and the **forcats** package

```
library(forcats)
```

419 Categorical variables can be stored in R as character strings in `character` or
 420 `factor` objects. A `factor` looks like a `character`, but it actually is an `integer`
 421 vector, where each `integer` is mapped to a `character` label. With this respect
 422 it is sort of an enhanced version of `character`. For example,

```
my_char_vec <- c("Pratik", "Theo", "Raph")
my_char_vec
#> [1] "Pratik" "Theo"  "Raph"
```

423 is a `character` vector, recognizable to its double quotes, while

```
my_fact_vec <- factor(my_char_vec) # as.factor would work too
my_fact_vec
#> [1] Pratik Theo  Raph
#> Levels: Pratik Raph Theo
```

424 is a `factor`, of which the *labels* are displayed. The *levels* of the factor are the
 425 unique values that appear in the vector. If I added an extra occurrence of my
 426 name:

```
factor(c(my_char_vec, "Raph"))
#> [1] Pratik Theo  Raph  Raph
#> Levels: Pratik Raph Theo
```

427 we would still have the the same levels. Note that the levels are returned as a
 428 `character` vector in alphabetical order by the `levels` function:

```
levels(my_fact_vec)
#> [1] "Pratik" "Raph"  "Theo"
```

429 Why does it matter? Well, most operations on categorical variables can be
 430 performed on `character` or `factor` objects, so it does not matter so much
 431 which one you use for your own data. However, some functions in R require
 432 you to provide categorical variables in one specific format, and others may even
 433 implicitly convert your variables. In `ggplot2` for example, character vectors
 434 are converted into factors by default. So, it is always good to remember the
 435 differences and what type your variables are.

436 But this is a tidyverse tutorial, so I would like to introduce here the package
 437 `forcats`, which offers tools to manipulate factors. First of all, most tools from
 438 `stringr` will work on factors. The `forcats` functions expand the string manip-
 439 ulation toolbox with factor-specific utilities. Similar in philosophy to `stringr`
 440 where functions started with `str_`, in `forcats` most functions start with `fct_`.

441 I see two main ways `forcats` can come handy in the kind of data most people
 442 deal with: playing with the order of the levels of a factor and playing with the
 443 levels themselves. We will show here a few examples, but the full breadth of
 444 factor manipulations can be found online or in the excellent `forcats` cheatsheet.

2.4.1 Change the order of the levels

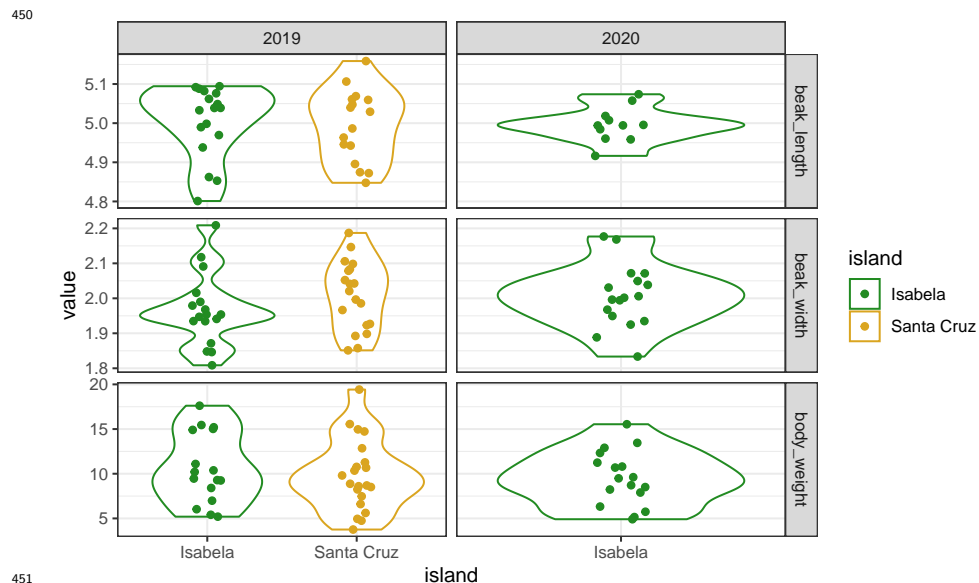
One example use-case where you would want to change the order of the levels of a factor is when plotting. Your categorical variable, for example, may not be plotted in the order you want. If we plot the distribution of each variable across islands, we get

```
# Make the plotting code a function so we can re-use it without copying and pasting
my_plot <- function(data) {

  # We do not cover the ggplot functions in this chapter, this is just to
  # illustrate our use-case, wait until chapter 5!
  library(ggplot2)
  ggplot(data, aes(x = island, y = value, color = island)) +
    geom_violin() +
    geom_jitter(width = 0.1) +
    facet_grid(variable ~ year, scales = "free") +
    theme_bw() +
    scale_color_manual(values = c("forestgreen", "goldenrod"))

}

my_plot(data)
# Remember that data are missing from Santa Cruz in 2020
```




Here, the islands (horizontal axis) and the variables (the facets) are displayed in alphabetical order. When making a figure you may want to customize these

orders in such a way that your message is optimally conveyed by your figure,
and this may involve playing with the order of levels.

Use `fct_relevel` to manually change the order of the levels:

```
data$island <- as.factor(data$island) # turn this column into a factor
data$island <- fct_relevel(data$island, c("Santa Cruz", "Isabela"))
my_plot(data) # order of islands has changed!
```



Beware that reordering a factor *does not change* the order of the items within the vector, only the order of the *levels*. So, it does not introduce any mismatch between the `island` column and the other columns! It only matters when the levels are called, for example, in a `ggplot`. As you can see:

```
data$island[1:10]
#> [1] Isabela Isabela Isabela Isabela Isabela Isabela Isabela Isabela Isabela Isabela
#> [10] Isabela
#> Levels: Santa Cruz Isabela
fct_relevel(data$island, c("Isabela", "Santa Cruz"))[1:10] # same thing, different levels
#> [1] Isabela Isabela Isabela Isabela Isabela Isabela Isabela Isabela Isabela Isabela
#> [10] Isabela
#> Levels: Isabela Santa Cruz
```

Alternatively, use `fct_inorder` to set the order of the levels to the order in which they appear:

```
data$variable <- as.factor(data$variable)
levels(data$variable)
#> [1] "beak_length" "beak_width" "body_weight"
```

```

levels(fct_inorder(data$variable))
#> [1] "body_weight" "beak_length" "beak_width"

```

465 or `fct_rev` to reverse the order of the levels:

```

levels(fct_rev(data$island)) # back in the alphabetical order
#> [1] "Isabela" "Santa Cruz"

```

466 Other variants exist to do more complex reordering, all present in the `forcats`
 467 cheatsheet, for example: * `fct_infreq` to re-order according to the frequency
 468 of each level (how many observation on each island?) * `fct_shift` to shift
 469 the order of all levels by a certain rank (in a circular way so that the last one
 470 becomes the first one or vice versa) * `fct_shuffle` if you want your levels in
 471 random order * `fct_reorder`, which reorders based on an associated variable
 472 (see `fct_reorder2` for even more complex relationship between the factor and
 473 the associated variable)

474 2.4.2 Change the levels themselves

475 Changing the levels of a factor will change the labels in the actual vector. It
 476 is similar to performing a string substitution in `stringr`. One can change the
 477 levels of a factor using `fct_recode`:

```

fct_recode(
  my_fact_vec,
  "Pratik Gupte" = "Pratik",
  "Theo Pannetier" = "Theo",
  "Raphael Scherrer" = "Raph"
)
#> [1] Pratik Gupte    Theo Pannetier    Raphael Scherrer
#> Levels: Pratik Gupte Raphael Scherrer Theo Pannetier

```

478 or collapse factor levels together using `fct_collapse`:

```

fct_collapse(my_fact_vec, EU = c("Theo", "Raph"), NonEU = "Pratik")
#> [1] NonEU EU      EU
#> Levels: NonEU EU

```

479 Again, we do not provide an exhaustive list of `forcats` functions here but the
 480 most usual ones, to give a glimpse of many things that one can do with factors.
 481 So, if you are dealing with factors, remember that `forcats` may have handy tools
 482 for you. Among others: * `fct_anon` to “anonymize”, i.e. replace the levels by
 483 random integers * `fct_lump` to collapse levels together based on their frequency
 484 (e.g. the two most frequent levels together)

2.4.3 Dropping levels

If you use factors in your tibble and get rid of one level, for any reason, the factor will usually remember the old levels, which may cause some problems when applying functions to your data.

```
data <- data[data$island == "Santa Cruz",] # keep only one island
unique(data$island) # Isabela is gone from the labels
#> [1] Santa Cruz
#> Levels: Santa Cruz Isabela
levels(data$island) # but not from the levels
#> [1] "Santa Cruz" "Isabela"
```

Use `droplevels` (from base R) to make sure you get rid of levels that are not in your data anymore:

```
data <- droplevels(data)
levels(data$island)
#> [1] "Santa Cruz"
```

Fortunately, most functions within the tidyverse will not complain about missing levels, and will automatically get rid of those inexistant levels for you. But because factors are such common causes of bugs, keep this in mind!

Note that this is equivalent to doing:

```
data$island <- fct_drop(data$island)
```

2.4.4 Other things

Among other things you can use in `forcats`: * `fct_count` to get the frequency of each level * `fct_c` to combine factors together

2.4.5 Take home message for forcats

Use this package to manipulate your factors. Do you need factors? Or are character vectors enough? That is your call, and may depend on the kind of analyses you want to do and what they require. We saw here that for plotting, having factors can allow you to do quite some tweaking of the display. If you encounter a situation where the order of encoding of your character vector starts to matter, then maybe converting into a factor would make your life easier. And if you do so, remember that lots of tools to perform all kinds of manipulation are available to you with both `stringr` and `forcats`.

507 2.5 External resources

508 Find lots of additional info by looking up the following links:

- 509 • The `readr/tibble/tidyr` and `forcats` cheatsheets.
- 510 • This link on the concept of tidy data
- 511 • The `tibble`, `tidyr` and `forcats` websites
- 512 • The `broom`, `tidymodels`, `tidygraph` and `tidytrees` websites

513 Chapter 3

514 Data manipulation with 515 dplyr

```
# load the tidyverse  
library(tidyverse)
```

516 3.1 Introduction

517 3.1.1 Foreword on dplyr

518 **dplyr** is tasked with performing all sorts of transformations on a dataset.

519 The structure of **dplyr** revolves around a set of functions, the so-called **verbs**,
520 that share a common syntax and logic, and are meant to work with one another
521 in chained operations. Chained operations are performed with the pipe operator
522 (`%>%`), that will be introduced in section 3.2.2.

523 The basic syntax is **verb(data, variable)**, where **data** is a data frame and
524 **variable** is the name of one or more columns containing a set of values for each
525 observation.

526 There are 5 main verbs, which names already hint at what they do: **rename()**,
527 **select()**, **filter()**, **mutate()**, and **summarise()**. I'm going to introduce each
528 of them (and a couple more) through the following sections.

529 3.1.2 Example data

530 Through this tutorial, we will be using mammal trait data from the Phylacine
531 database. Let's have a peek at what it contains.

```

phylacine <- read_csv("data/phylacine_traits.csv")
phylacine
#> # A tibble: 5,831 x 24
#>   Binomial.1.2 Order.1.2 Family.1.2 Genus.1.2 Species.1.2 Terrestrial Marine
#>   <chr>           <chr>      <chr>      <chr>      <chr>           <dbl> <dbl>
#> 1 Abditomys_l~ Rodentia Muridae Abditomys latidens           1      0
#> 2 Abeomelomys~ Rodentia Muridae Abeomelo~ sevia           1      0
#> 3 Abrawayaomy~ Rodentia Cricetidae Abrawaya~ ruschii           1      0
#> 4 Abrocoma_be~ Rodentia Abrocomid~ Abrocoma bennettii           1      0
#> 5 Abrocoma_bo~ Rodentia Abrocomid~ Abrocoma boliviensis           1      0
#> 6 Abrocoma_bu~ Rodentia Abrocomid~ Abrocoma budini           1      0
#> # ... with 5,825 more rows, and 17 more variables: Freshwater <dbl>,
#> # Aerial <dbl>, Life.Habit.Method <chr>, Life.Habit.Source <chr>,
#> # Mass.g <dbl>, Mass.Method <chr>, Mass.Source <chr>, Mass.Comparison <chr>,
#> # Mass.Comparison.Source <chr>, Island.Endemicity <chr>,
#> # IUCN.Status.1.2 <chr>, Added.IUCN.Status.1.2 <chr>, Diet.Plant <dbl>,
#> # Diet.Vertebrate <dbl>, Diet.Invertebrate <dbl>, Diet.Method <chr>,
#> # Diet.Source <chr>

```

532 `readr` automatically loads the data in a `tibble`, as we have seen in chapter 1
 533 and 2. Calling the `tibble` gives a nice preview of what it contains. We have data
 534 for 5,831 mammal species, and the variables contain information on taxonomy,
 535 (broad) habitat, mass, IUCN status, and diet.

536 If you remember Section 1.2 on tidy data, you may see that this data isn't
 537 exactly tidy. In fact, some columns are in wide (and messy) format, like the
 538 "habitat" (terrestrial, marine, etc.) and diet columns.

539 `dplyr` actually does not require your data to be strictly tidy. If you feel that
 540 your data satisfies the definition "one observation per row, one variable per
 541 column", that's probably good enough.

542 I use a `tibble` here, but `dplyr` works equally well on base data frames. In fact,
 543 `dplyr` is built for `data.frame` objects, and tibbles are data frames. Therefore,
 544 tibbles are mortal.

545 3.2 Working with existing variables

546 3.2.1 Renaming variables with `rename()`

547 The variable names in the `phylacine` dataset are descriptive, but quite unpracti-
 548 cal. Typing `Binomial.1.2` is cumbersome and subject to typos (in fact, I just
 549 made one). `binomial` would be much simpler to use.

550 Changing names is straightforward with `rename()`.

```

rename(.data = phylacine, "binomial" = Binomial.1.2)
#> # A tibble: 5,831 x 24
#>   binomial Order.1.2 Family.1.2 Genus.1.2 Species.1.2 Terrestrial Marine
#>   <chr>      <chr>      <chr>      <chr>      <chr>      <dbl>  <dbl>
#> 1 Abditom~ Rodentia Muridae  Abditomys latidens      1      0
#> 2 Abeomel~ Rodentia Muridae  Abeomelo~ sevia      1      0
#> 3 Abraway~ Rodentia Cricetidae Abrawaya~ ruschii      1      0
#> 4 Abrocom~ Rodentia Abrocomid~ Abrocoma bennettii      1      0
#> 5 Abrocom~ Rodentia Abrocomid~ Abrocoma boliviensis      1      0
#> 6 Abrocom~ Rodentia Abrocomid~ Abrocoma budini      1      0
#> # ... with 5,825 more rows, and 17 more variables: Freshwater <dbl>,
#> #   Aerial <dbl>, Life.Habit.Method <chr>, Life.Habit.Source <chr>,
#> #   Mass.g <dbl>, Mass.Method <chr>, Mass.Source <chr>, Mass.Comparison <chr>,
#> #   Mass.Comparison.Source <chr>, Island.Endemicity <chr>,
#> #   IUCN.Status.1.2 <chr>, Added.IUCN.Status.1.2 <chr>, Diet.Plant <dbl>,
#> #   Diet.Vertebrate <dbl>, Diet.Invertebrate <dbl>, Diet.Method <chr>,
#> #   Diet.Source <chr>

```

551 The first argument is always `.data`, the data table you want to apply change
 552 to. Note how columns are referred to. Once the data table has been passed as an
 553 argument, there is no need to refer to it directly anymore, `dplyr` understands
 554 that you're dealing with variables inside that data frame. So drop that `data$var`,
 555 `data[, "var"]`, and forget the very existence of `attach()` / `detach()`.

556 You can refer to variables names either with strings or directly as objects,
 557 whether you're reading or creating them:

```

rename(
  phylacine,
  # this works
  binomial = Binomial.1.2
)
rename(
  phylacine,
  # this works too!
  binomial = "Binomial.1.2"
)
rename(
  phylacine,
  # guess what
  "binomial" = "Binomial.1.2"
)

```

558 I have applied similar changes to all variables in the dataset. Here is what the
 559 new names look like:

```

560 #> # A tibble: 5,831 x 24
561 #>   binomial order family genus species terrestrial marine freshwater aerial

```

```

562 #>   <chr>      <chr> <chr> <chr> <chr>          <dbl> <dbl>      <dbl> <dbl>
563 #> 1 Abditom~ Rode~ Murid~ Abdi~ latide~          1      0          0      0
564 #> 2 Abeomel~ Rode~ Murid~ Abeo~ sevia          1      0          0      0
565 #> 3 Abraway~ Rode~ Crice~ Abra~ ruschii          1      0          0      0
566 #> 4 Abrocom~ Rode~ Abroc~ Abro~ bennet~          1      0          0      0
567 #> 5 Abrocom~ Rode~ Abroc~ Abro~ bolivi~          1      0          0      0
568 #> 6 Abrocom~ Rode~ Abroc~ Abro~ budini          1      0          0      0
569 #> # ... with 5,825 more rows, and 15 more variables: life_habit_method <chr>,
570 #> #   life_habit_source <chr>, mass_g <dbl>, mass_method <chr>,
571 #> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
572 #> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>,
573 #> #   diet_plant <dbl>, diet_vertibrate <dbl>, diet_invertebrate <dbl>,
574 #> #   diet_method <chr>, diet_source <chr>

```

3.2.2 The pipe operator %>%

If you have already come across pieces of code using the tidyverse, chances are that you have seen this odd symbol. While the pipe is not strictly-speaking a part of the tidyverse (it comes from its own package, `magrittr`), it is imported along with each package and widely used in conjunction with its functions. What does it do? Consider the following example with `rename()`:

```

phylacine2 <- readr::read_csv("data/phylacine_traits.csv")
# regular syntax
rename(phylacine2, "binomial" = "Binomial.1.2")
#> # A tibble: 5,831 x 24
#>   binomial Order.1.2 Family.1.2 Genus.1.2 Species.1.2 Terrestrial Marine
#>   <chr>      <chr>      <chr>      <chr>      <chr>          <dbl> <dbl>
#> 1 Abditom~ Rodentia Muridae   Abditomys latidens          1      0
#> 2 Abeomel~ Rodentia Muridae   Abeomelo~ sevia          1      0
#> 3 Abraway~ Rodentia Cricetidae Abrawaya~ ruschii          1      0
#> 4 Abrocom~ Rodentia Abrocomid~ Abrocoma bennettii          1      0
#> 5 Abrocom~ Rodentia Abrocomid~ Abrocoma boliviensis          1      0
#> 6 Abrocom~ Rodentia Abrocomid~ Abrocoma budini          1      0
#> # ... with 5,825 more rows, and 17 more variables: Freshwater <dbl>,
#> #   Aerial <dbl>, Life.Habit.Method <chr>, Life.Habit.Source <chr>,
#> #   Mass.g <dbl>, Mass.Method <chr>, Mass.Source <chr>, Mass.Comparison <chr>,
#> #   Mass.Comparison.Source <chr>, Island.Endemicity <chr>,
#> #   IUCN.Status.1.2 <chr>, Added.IUCN.Status.1.2 <chr>, Diet.Plant <dbl>,
#> #   Diet.Vertibrate <dbl>, Diet.Invertebrate <dbl>, Diet.Method <chr>,
#> #   Diet.Source <chr>
# alternative syntax with the pipe operator
phylacine2 %>% rename("binomial" = "Binomial.1.2")
#> # A tibble: 5,831 x 24
#>   binomial Order.1.2 Family.1.2 Genus.1.2 Species.1.2 Terrestrial Marine

```

```

#>   <chr>   <chr>   <chr>   <chr>   <chr>   <dbl> <dbl>
#> 1 Abditom~ Rodentia Muridae Abditomys latidens      1      0
#> 2 Abeomel~ Rodentia Muridae Abeomelo~ sevia      1      0
#> 3 Abraway~ Rodentia Cricetidae Abrawaya~ ruschii      1      0
#> 4 Abrocom~ Rodentia Abrocomid~ Abrocoma bennettii      1      0
#> 5 Abrocom~ Rodentia Abrocomid~ Abrocoma boliviensis      1      0
#> 6 Abrocom~ Rodentia Abrocomid~ Abrocoma budini      1      0
#> # ... with 5,825 more rows, and 17 more variables: Freshwater <dbl>,
#> #   Aerial <dbl>, Life.Habit.Method <chr>, Life.Habit.Source <chr>,
#> #   Mass.g <dbl>, Mass.Method <chr>, Mass.Source <chr>, Mass.Comparison <chr>,
#> #   Mass.Comparison.Source <chr>, Island.Endemicity <chr>,
#> #   IUCN.Status.1.2 <chr>, Added.IUCN.Status.1.2 <chr>, Diet.Plant <dbl>,
#> #   Diet.Vertebrate <dbl>, Diet.Invertebrate <dbl>, Diet.Method <chr>,
#> #   Diet.Source <chr>

```

581 Got it? The pipe takes the object on its left-side and silently feeds it to the
 582 *first* argument of the function on its right-side. It could be read as “take x,
 583 then do...”. The reason for using the pipe is because it makes code syntax
 584 closer to the syntax of a sentence, and therefore, easier and faster for your brain
 585 to process (and write!) the code. In particular, the pipe enables easy chains
 586 of operations, where you apply something to an object, then apply something
 587 else to the outcome, and so on... Through the later sections, you will see some
 588 examples of chained operations with `dplyr` functions, but for that I first need
 589 to introduce a couple more verbs.

590 Using the pipe can be quite unsettling at first, because you are not used to
 591 think in this way. But if you push a bit for it, I promise it will make things
 592 a lot easier (and it’s quite addictive!). To avoid typing the tedious symbols,
 593 `magrittr` installs a shortcut for you in RStudio. Use **Ctrl + Shift + M** on
 594 Windows, and **Cmd + Shift + M** on MacOS.

595 Finally I should emphasize that the use of the pipe isn’t limited to the tidyverse,
 596 but extends to almost all R functions. Remember that by default the piped value
 597 is always matched to the first argument of the following function

```

5  %>% rep(3)
#> [1] 5 5 5
"meow" %>% cat()
#> meow

```

598 If you need to pass the left-hand side to an argument other than the first, you
 599 can use the dot place-holder `..`

```

"meow" %>% cat("cats", "go")
#> meow cats go

```

600 Because of its syntax, most base R operators are not compatible with the pipe
 601 (but this is very rarely needed). If needed, `magrittr` introduces alternative
 602 functions for operators.

603 Subsetting operators can be piped, with the dot place-holder.

```
# 5 %>% * 3 # no, won't work
# 5 %>% .* 3 # neither
5 %>% magrittr::multiply_by(3) # yes
#> [1] 15

# subsetting
list("monkey see", "monkey_do") %>% .[[2]]
#> [1] "monkey_do"
phylacine %>% .$binomial %>% head()
#> [1] "Abditomys_latidens" "Abeomelomys_sevia" "Abrawayaomys_ruschii"
#> [4] "Abrocoma_bennettii" "Abrocoma_boliviensis" "Abrocoma_budini"
```

604 Because subsetting in this way is particularly hideous, **dplyr** delivers a function
605 to extract values from a single variable. It only works on tables, though.

```
phylacine %>% pull(binomial) %>% head()
#> [1] "Abditomys_latidens" "Abeomelomys_sevia" "Abrawayaomys_ruschii"
#> [4] "Abrocoma_bennettii" "Abrocoma_boliviensis" "Abrocoma_budini"
```

606 3.2.3 Select variables with **select()**

607 To extract a set of variables (i.e. columns), use the conveniently-named
608 **select()**. The basic syntax is the same as **rename()**: pass your data as the
609 first argument, then call the variables to select, quoted or not.

```
# Single variable
phylacine %>% select(binomial)
#> # A tibble: 5,831 x 1
#>   binomial
#>   <chr>
#> 1 Abditomys_latidens
#> 2 Abeomelomys_sevia
#> 3 Abrawayaomys_ruschii
#> 4 Abrocoma_bennettii
#> 5 Abrocoma_boliviensis
#> 6 Abrocoma_budini
#> # ... with 5,825 more rows
# A set of variables
phylacine %>% select(genus, "species", mass_g)
#> # A tibble: 5,831 x 3
#>   genus      species      mass_g
#>   <chr>      <chr>      <dbl>
#> 1 Abditomys  latidens      269
#> 2 Abeomelomys sevia      52
#> 3 Abrawayaomys ruschii      63
```



```
#> 4 Abrocoma      bennettii      250
#> 5 Abrocoma      boliviensis    158
#> 6 Abrocoma      budini         361.
#> # ... with 5,825 more rows
# A range of contiguous variables
phylacine %>% select(family:terrestrial)
#> # A tibble: 5,831 x 4
#>   family      genus      species      terrestrial
#>   <chr>      <chr>      <chr>      <dbl>
#> 1 Muridae    Abditomys    latidens      1
#> 2 Muridae    Abeomelomys sevia      1
#> 3 Cricetidae Abrawayaomys ruschii      1
#> 4 Abrocomidae Abrocoma      bennettii      1
#> 5 Abrocomidae Abrocoma      boliviensis    1
#> 6 Abrocomidae Abrocoma      budini         1
#> # ... with 5,825 more rows
```

610 You can select by variable numbers. This is not recommended, as prone to
 611 errors, especially if you change the variable order.

```
phylacine %>% select(2)
#> # A tibble: 5,831 x 1
#>   order
#>   <chr>
#> 1 Rodentia
#> 2 Rodentia
#> 3 Rodentia
#> 4 Rodentia
#> 5 Rodentia
#> 6 Rodentia
#> # ... with 5,825 more rows
```

612 `select()` can also be used to *exclude* variables:

```
phylacine %>% select(-binomial)
#> # A tibble: 5,831 x 23
#>   order family genus species terrestrial marine freshwater aerial
#>   <chr> <chr> <chr> <chr>      <dbl> <dbl>      <dbl> <dbl>
#> 1 Rode~ Murid~ Abdi~ latide~      1     0         0     0
#> 2 Rode~ Murid~ Abeo~ sevia      1     0         0     0
#> 3 Rode~ Crice~ Abra~ ruschii      1     0         0     0
#> 4 Rode~ Abroc~ Abro~ bennet~      1     0         0     0
#> 5 Rode~ Abroc~ Abro~ bolivi~      1     0         0     0
#> 6 Rode~ Abroc~ Abro~ budini      1     0         0     0
#> # ... with 5,825 more rows, and 15 more variables: life_habit_method <chr>,
#> #   life_habit_source <chr>, mass_g <dbl>, mass_method <chr>,
#> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
#> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>,
```

```
#> #   diet_plant <dbl>, diet_vertibrate <dbl>, diet_invertebrate <dbl>,
#> #   diet_method <chr>, diet_source <chr>
phylacine %>% select(-(binomial:species))
#> # A tibble: 5,831 x 19
#>   terrestrial marine freshwater aerial life_habit_meth~ life_habit_sour~ mass_g
#>   <dbl> <dbl> <dbl> <dbl> <chr> <chr> <dbl>
#> 1      1      0      0      0 0 Reported IUCN. 2016. IUC~ 269
#> 2      1      0      0      0 0 Reported IUCN. 2016. IUC~ 52
#> 3      1      0      0      0 0 Reported IUCN. 2016. IUC~ 63
#> 4      1      0      0      0 0 Reported IUCN. 2016. IUC~ 250
#> 5      1      0      0      0 0 Reported IUCN. 2016. IUC~ 158
#> 6      1      0      0      0 0 Reported IUCN. 2016. IUC~ 361.
#> # ... with 5,825 more rows, and 12 more variables: mass_method <chr>,
#> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
#> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>,
#> #   diet_plant <dbl>, diet_vertibrate <dbl>, diet_invertebrate <dbl>,
#> #   diet_method <chr>, diet_source <chr>
```

613 `select()` and `rename()` are pretty similar, and in fact, `select()` can also
 614 rename variables along the way:

```
phylacine %>% select("linnaeus" = binomial)
#> # A tibble: 5,831 x 1
#>   linnaeus
#>   <chr>
#> 1 Abditomys_latidens
#> 2 Abeomelomys_sevia
#> 3 Abrawayaomys_ruschii
#> 4 Abrocoma_bennettii
#> 5 Abrocoma_boliviensis
#> 6 Abrocoma_budini
#> # ... with 5,825 more rows
```

615 And you can mix all of that at once:

```
phylacine %>% select(
  "fam" = family,
  genus:freshwater,
  -terrestrial
)
#> # A tibble: 5,831 x 5
#>   fam      genus      species      marine freshwater
#>   <chr>    <chr>    <chr>    <dbl>    <dbl>
#> 1 Muridae Abditomys latidens      0      0
#> 2 Muridae Abeomelomys sevia      0      0
#> 3 Cricetidae Abrawayaomys ruschii      0      0
#> 4 Abrocomidae Abrocoma bennettii      0      0
#> 5 Abrocomidae Abrocoma boliviensis      0      0
```

```
#> 6 Abrocomidae Abrocoma budini 0 0
#> # ... with 5,825 more rows
```

616 3.2.4 Select variables with helpers

617 The Rstudio team just released `dplyr` 1.0.0, and along with it, some nice
 618 helper functions to ease the selection of a set of variables. I give three examples
 619 here, and encourage you to look at the documentation (`?select()`) to find out
 620 more.

```
phylacine %>% select(where(is.numeric))
#> # A tibble: 5,831 x 8
#>   terrestrial marine freshwater aerial mass_g diet_plant diet_vertbrate
#>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1      1      0      0      0 269      100      0
#> 2      1      0      0      0 52       78       3
#> 3      1      0      0      0 63       88       1
#> 4      1      0      0      0 250      100      0
#> 5      1      0      0      0 158      100      0
#> 6      1      0      0      0 361.     100      0
#> # ... with 5,825 more rows, and 1 more variable: diet_invertebrate <dbl>
phylacine %>% select(contains("mass") | contains("diet"))
#> # A tibble: 5,831 x 10
#>   mass_g mass_method mass_source mass_comparison mass_comparison~ diet_plant
#>   <dbl> <chr> <chr> <chr> <chr> <dbl>
#> 1 269 Reported Smith, F. ~ <NA> <NA> 100
#> 2 52 Reported Smith, F. ~ <NA> <NA> 78
#> 3 63 Reported Smith, F. ~ <NA> <NA> 88
#> 4 250 Reported Smith, F. ~ <NA> <NA> 100
#> 5 158 Reported Smith, F. ~ <NA> <NA> 100
#> 6 361. Assumed is~ Journal of~ Abrocoma_ciner~ Journal of Mamm~ 100
#> # ... with 5,825 more rows, and 4 more variables: diet_vertbrate <dbl>,
#> # diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>
```

```
habitats <- c("terrestrial", "marine", "arboreal", "fossorial", "freshwater")
phylacine %>% select(any_of(habitats))
#> # A tibble: 5,831 x 3
#>   terrestrial marine freshwater
#>   <dbl> <dbl> <dbl>
#> 1      1      0      0
#> 2      1      0      0
#> 3      1      0      0
#> 4      1      0      0
#> 5      1      0      0
#> 6      1      0      0
```

```
#> # ... with 5,825 more rows
```

3.2.5 Rearranging variable order with relocate()

The order of variables seldom matters in `dplyr`, but due to popular demand, `dplyr` now has a dedicated verb to rearrange the order of variables. The syntax is identical to `rename()`, `select()`.

```
phylacine %>% relocate(mass_g, .before = binomial)
#> # A tibble: 5,831 x 24
#>   mass_g binomial order family genus species terrestrial marine freshwater
#>   <dbl> <chr>    <chr> <chr> <chr> <chr>          <dbl>  <dbl>      <dbl>
#> 1   269  Abditom~ Rode~ Murid~ Abdi~ latide~         1      0          0
#> 2    52  Abeomel~ Rode~ Murid~ Abeo~ sevia         1      0          0
#> 3    63  Abraway~ Rode~ Crice~ Abra~ ruschii       1      0          0
#> 4   250  Abrocom~ Rode~ Abroc~ Abro~ bennet~       1      0          0
#> 5   158  Abrocom~ Rode~ Abroc~ Abro~ bolivi~       1      0          0
#> 6  361. Abrocom~ Rode~ Abroc~ Abro~ budini       1      0          0
#> # ... with 5,825 more rows, and 15 more variables: aerial <dbl>,
#> #   life_habit_method <chr>, life_habit_source <chr>, mass_method <chr>,
#> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
#> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>,
#> #   diet_plant <dbl>, diet_vertibrate <dbl>, diet_invertebrate <dbl>,
#> #   diet_method <chr>, diet_source <chr>
phylacine %>% relocate(starts_with("diet"), .after = species)
#> # A tibble: 5,831 x 24
#>   binomial order family genus species diet_plant diet_vertibrate
#>   <chr>    <chr> <chr> <chr> <chr>          <dbl>          <dbl>
#> 1 Abditom~ Rode~ Murid~ Abdi~ latide~        100            0
#> 2 Abeomel~ Rode~ Murid~ Abeo~ sevia           78            3
#> 3 Abraway~ Rode~ Crice~ Abra~ ruschii          88            1
#> 4 Abrocom~ Rode~ Abroc~ Abro~ bennet~        100            0
#> 5 Abrocom~ Rode~ Abroc~ Abro~ bolivi~        100            0
#> 6 Abrocom~ Rode~ Abroc~ Abro~ budini          100            0
#> # ... with 5,825 more rows, and 17 more variables: diet_invertebrate <dbl>,
#> #   diet_method <chr>, diet_source <chr>, terrestrial <dbl>, marine <dbl>,
#> #   freshwater <dbl>, aerial <dbl>, life_habit_method <chr>,
#> #   life_habit_source <chr>, mass_g <dbl>, mass_method <chr>,
#> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
#> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>
```

625 3.3 Working with observations

626 3.3.1 Ordering rows by value - arrange()

627 `arrange()` sorts rows in the data by **ascending** value for a given variable. Use
628 the wrapper `desc()` to sort by descending values instead.

```
# Smallest mammals
phylacine %>%
  arrange(mass_g) %>%
  select(binomial, mass_g)
#> # A tibble: 5,831 x 2
#>   binomial      mass_g
#>   <chr>         <dbl>
#> 1 Sorex_yukonicus    1.6
#> 2 Crocidura_levicula  1.8
#> 3 Suncus_remyi      1.8
#> 4 Crocidura_lusitania 2
#> 5 Kerivoula_minuta  2.1
#> 6 Suncus_etruscus   2.1
#> # ... with 5,825 more rows

# Largest mammals
phylacine %>%
  arrange(desc(mass_g)) %>%
  select(binomial, mass_g)
#> # A tibble: 5,831 x 2
#>   binomial      mass_g
#>   <chr>         <dbl>
#> 1 Balaenoptera_musculus 190000000
#> 2 Balaena_mysticetus   100000000
#> 3 Balaenoptera_physalus 70000000
#> 4 Caperea_marginata    32000000
#> 5 Megaptera_novaeangliae 30000000
#> 6 Eschrichtius_robustus 28500000
#> # ... with 5,825 more rows

# Extra variables are used to sort ties in the first variable
phylacine %>%
  arrange(mass_g, desc(binomial)) %>%
  select(binomial, mass_g)
#> # A tibble: 5,831 x 2
#>   binomial      mass_g
#>   <chr>         <dbl>
#> 1 Sorex_yukonicus    1.6
```

```
#> 2 Suncus_remyi          1.8
#> 3 Crocidura_levicula    1.8
#> 4 Crocidura_lusitania    2
#> 5 Suncus_etruscus       2.1
#> 6 Kerivoula_minuta      2.1
#> # ... with 5,825 more rows
```

629 *Important:* NA values, if present, are always ordered at the end!

630 3.3.2 Subset rows by position - `slice()`

631 Use `slice()` and its variants to extract particular rows.

```
phylacine %>% slice(3) # third row
#> # A tibble: 1 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl> <dbl>          <dbl> <dbl>
#> 1 Abraway~ Rode~ Crice~ Abra~ ruschii          1      0              0      0
#> # ... with 15 more variables: life_habit_method <chr>, life_habit_source <chr>,
#> #   mass_g <dbl>, mass_method <chr>, mass_source <chr>, mass_comparison <chr>,
#> #   mass_comparison_source <chr>, island_endemicity <chr>, iucn_status <chr>,
#> #   added_iucn_status <chr>, diet_plant <dbl>, diet_vertibrate <dbl>,
#> #   diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>
phylacine %>% slice(5, 1, 2) # fifth, first and second row
#> # A tibble: 3 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl> <dbl>          <dbl> <dbl>
#> 1 Abrocom~ Rode~ Abroc~ Abro~ bolivi~          1      0              0      0
#> 2 Abditom~ Rode~ Murid~ Abdi~ latide~          1      0              0      0
#> 3 Abeomel~ Rode~ Murid~ Abeo~ sevia          1      0              0      0
#> # ... with 15 more variables: life_habit_method <chr>, life_habit_source <chr>,
#> #   mass_g <dbl>, mass_method <chr>, mass_source <chr>, mass_comparison <chr>,
#> #   mass_comparison_source <chr>, island_endemicity <chr>, iucn_status <chr>,
#> #   added_iucn_status <chr>, diet_plant <dbl>, diet_vertibrate <dbl>,
#> #   diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>
phylacine %>% slice(rep(3, 2)) # duplicate the third row
#> # A tibble: 2 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl> <dbl>          <dbl> <dbl>
#> 1 Abraway~ Rode~ Crice~ Abra~ ruschii          1      0              0      0
#> 2 Abraway~ Rode~ Crice~ Abra~ ruschii          1      0              0      0
#> # ... with 15 more variables: life_habit_method <chr>, life_habit_source <chr>,
#> #   mass_g <dbl>, mass_method <chr>, mass_source <chr>, mass_comparison <chr>,
#> #   mass_comparison_source <chr>, island_endemicity <chr>, iucn_status <chr>,
#> #   added_iucn_status <chr>, diet_plant <dbl>, diet_vertibrate <dbl>,
#> #   diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>
```

```

phylacine %>% slice(-c(2:5830)) # exclude all but first and last row
#> # A tibble: 2 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl>  <dbl>          <dbl>  <dbl>
#> 1 Abditom~ Rode~ Murid~ Abdi~ latide~          1    0              0    0
#> 2 Zyzomys~ Rode~ Murid~ Zyzo~ woodwa~          1    0              0    0
#> # ... with 15 more variables: life_habit_method <chr>, life_habit_source <chr>,
#> #   mass_g <dbl>, mass_method <chr>, mass_source <chr>, mass_comparison <chr>,
#> #   mass_comparison_source <chr>, island_endemicity <chr>, iucn_status <chr>,
#> #   added_iucn_status <chr>, diet_plant <dbl>, diet_vertibrate <dbl>,
#> #   diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>

phylacine %>% slice_tail(n = 3) # last three rows
#> # A tibble: 3 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl>  <dbl>          <dbl>  <dbl>
#> 1 Zyzomys~ Rode~ Murid~ Zyzo~ palata~          1    0              0    0
#> 2 Zyzomys~ Rode~ Murid~ Zyzo~ pedunc~          1    0              0    0
#> 3 Zyzomys~ Rode~ Murid~ Zyzo~ woodwa~          1    0              0    0
#> # ... with 15 more variables: life_habit_method <chr>, life_habit_source <chr>,
#> #   mass_g <dbl>, mass_method <chr>, mass_source <chr>, mass_comparison <chr>,
#> #   mass_comparison_source <chr>, island_endemicity <chr>, iucn_status <chr>,
#> #   added_iucn_status <chr>, diet_plant <dbl>, diet_vertibrate <dbl>,
#> #   diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>

phylacine %>% slice_max(mass_g) # largest mammal
#> # A tibble: 1 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl>  <dbl>          <dbl>  <dbl>
#> 1 Balaeno~ Ceta~ Balae~ Bala~ muscul~          0    1              0    0
#> # ... with 15 more variables: life_habit_method <chr>, life_habit_source <chr>,
#> #   mass_g <dbl>, mass_method <chr>, mass_source <chr>, mass_comparison <chr>,
#> #   mass_comparison_source <chr>, island_endemicity <chr>, iucn_status <chr>,
#> #   added_iucn_status <chr>, diet_plant <dbl>, diet_vertibrate <dbl>,
#> #   diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>

```

632 You can also sample random rows in the data:

```

phylacine %>% slice_sample() # a random row
#> # A tibble: 1 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl>  <dbl>          <dbl>  <dbl>
#> 1 Crocidu~ Euli~ Soric~ Croc~ levicu~          1    0              0    0
#> # ... with 15 more variables: life_habit_method <chr>, life_habit_source <chr>,
#> #   mass_g <dbl>, mass_method <chr>, mass_source <chr>, mass_comparison <chr>,
#> #   mass_comparison_source <chr>, island_endemicity <chr>, iucn_status <chr>,
#> #   added_iucn_status <chr>, diet_plant <dbl>, diet_vertibrate <dbl>,
#> #   diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>

```

```

# bootstrap
phylacine %>% slice_sample(n = 5831, replace = TRUE)
#> # A tibble: 5,831 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl>  <dbl>          <dbl>  <dbl>
#> 1 Rhinolo~ Chir~ Rhino~ Rhin~ adami          0      0              0      1
#> 2 Hylomys~ Euli~ Erina~ Hylo~ megal~          1      0              0      0
#> 3 Sciurus~ Rode~ Sciur~ Sciu~ yucata~          1      0              0      0
#> 4 Emballo~ Chir~ Embal~ Emba~ alecto          0      0              0      1
#> 5 Pteralo~ Chir~ Ptero~ Pter~ taki          0      0              0      1
#> 6 Lasiorh~ Dipr~ Vomba~ Lasi~ latifr~          1      0              0      0
#> # ... with 5,825 more rows, and 15 more variables: life_habit_method <chr>,
#> #   life_habit_source <chr>, mass_g <dbl>, mass_method <chr>,
#> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
#> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>,
#> #   diet_plant <dbl>, diet_vertibrate <dbl>, diet_invertebrate <dbl>,
#> #   diet_method <chr>, diet_source <chr>

```

633 3.3.3 Subsetting rows by value with filter()

634 `filter()` does a similar job as `slice()`, but extract rows that satisfy a set of
 635 conditions. The conditions are supplied much the same way as you would do
 636 for an `if` statement.

637 Along with `mutate()` (next section), this is probably the function you are going
 638 to use the most.

639 For example, I might want to extract mammals above a given mass:

```

# megafauna
phylacine %>%
  filter(mass_g > 1e5) %>% # 100 kg
  select(binomial, mass_g)
#> # A tibble: 302 x 2
#>   binomial          mass_g
#>   <chr>          <dbl>
#> 1 Ailuropoda_melanoleuca 108400
#> 2 Alcelaphus_buselaphus 171002.
#> 3 Alces_alces          356998
#> 4 Archaeoindris_fontoynonti 160000
#> 5 Arctocephalus_forsteri 101250
#> 6 Arctocephalus_pusillus 178500
#> # ... with 296 more rows

# non-extinct megafauna

```



```

phylacine %>%
  filter(mass_g > 1e5, iucn_status != "EP") %>%
  select(binomial, mass_g, iucn_status)
#> # A tibble: 178 x 3
#>   binomial      mass_g iucn_status
#>   <chr>          <dbl> <chr>
#> 1 Ailuropoda_melanoleuca 108400 VU
#> 2 Alcelaphus_buselaphus 171002 LC
#> 3 Alces_alces          356998 LC
#> 4 Arctocephalus_forsteri 101250 LC
#> 5 Arctocephalus_pusillus 178500 LC
#> 6 Arctocephalus_townsendi 105000 LC
#> # ... with 172 more rows

```

640 Are there any flying mammals that aren't bats?

```

phylacine %>%
  filter(aerial == 1, order != "Chiroptera")
#> # A tibble: 0 x 24
#> # ... with 24 variables: binomial <chr>, order <chr>, family <chr>,
#> #   genus <chr>, species <chr>, terrestrial <dbl>, marine <dbl>,
#> #   freshwater <dbl>, aerial <dbl>, life_habit_method <chr>,
#> #   life_habit_source <chr>, mass_g <dbl>, mass_method <chr>,
#> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
#> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>,
#> #   diet_plant <dbl>, diet_vertibrate <dbl>, diet_invertebrate <dbl>,
#> #   diet_method <chr>, diet_source <chr>
# no :(

```

641 Are humans included in the table?

```

phylacine %>% filter(binomial == "Homo_sapiens")
#> # A tibble: 1 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl> <dbl>          <dbl> <dbl>
#> 1 Homo_sa~ Prim~ Homin~ Homo sapiens          1      0              0      0
#> # ... with 15 more variables: life_habit_method <chr>, life_habit_source <chr>,
#> #   mass_g <dbl>, mass_method <chr>, mass_source <chr>, mass_comparison <chr>,
#> #   mass_comparison_source <chr>, island_endemicity <chr>, iucn_status <chr>,
#> #   added_iucn_status <chr>, diet_plant <dbl>, diet_vertibrate <dbl>,
#> #   diet_invertebrate <dbl>, diet_method <chr>, diet_source <chr>

```

642 filter() can be used to deal with NAs:

```

phylacine %>%
  filter(!is.na(mass_comparison))
#> # A tibble: 754 x 24
#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl> <dbl>          <dbl> <dbl>

```

```
#> 1 Abrocom~ Rode~ Abroc~ Abro~ budini          1      0      0      0
#> 2 Abrocom~ Rode~ Abroc~ Abro~ famati~          1      0      0      0
#> 3 Abrocom~ Rode~ Abroc~ Abro~ shista~          1      0      0      0
#> 4 Abrocom~ Rode~ Abroc~ Abro~ uspoll~          1      0      0      0
#> 5 Abrocom~ Rode~ Abroc~ Abro~ vaccar~          1      0      0      0
#> 6 Acerodo~ Chir~ Ptero~ Acer~ humilis          0      0      0      1
#> # ... with 748 more rows, and 15 more variables: life_habit_method <chr>,
#> #   life_habit_source <chr>, mass_g <dbl>, mass_method <chr>,
#> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
#> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>,
#> #   diet_plant <dbl>, diet_vertibrate <dbl>, diet_invertebrate <dbl>,
#> #   diet_method <chr>, diet_source <chr>
```

643 Tip: `dplyr` introduces the useful function `between()` that does exactly what
 644 the name implies

```
between(1:5, 2, 4)
#> [1] FALSE TRUE TRUE TRUE FALSE

# Mesofauna
phylacine %>%
  filter(mass_g > 1e3, mass_g < 1e5) %>%
  select(binomial, mass_g)
#> # A tibble: 1,126 x 2
#>   binomial      mass_g
#>   <chr>         <dbl>
#> 1 Acerodon_jubatus    1075
#> 2 Acinonyx_jubatus   46700
#> 3 Acratocnus_odontrigonus 22990
#> 4 Acratocnus_ye      21310
#> 5 Addax_nasomaculatus  70000.
#> 6 Aepyceros_melampus   52500.
#> # ... with 1,120 more rows

# same thing
phylacine %>%
  filter(mass_g %>% between(1e3, 1e5)) %>%
  select(binomial, mass_g)
#> # A tibble: 1,148 x 2
#>   binomial      mass_g
#>   <chr>         <dbl>
#> 1 Acerodon_jubatus    1075
#> 2 Acinonyx_jubatus   46700
#> 3 Acratocnus_odontrigonus 22990
#> 4 Acratocnus_ye      21310
#> 5 Addax_nasomaculatus  70000.
#> 6 Aepyceros_melampus   52500.
```

```
#> # ... with 1,142 more rows
```

645 Note that you can pipe operations inside function arguments as in the last line
646 above (arguments are expressions, after all!).

647 3.4 Making new variables

648 3.4.1 Create new variables with `mutate()`

649 Very often in data analysis, you will want to create new variables, or edit existing
650 ones. This is done easily through `mutate()`. For example, consider the diet data:

```
diet <- phylacine %>%
  select(
    binomial,
    contains("diet") & !contains(c("method", "source"))
  )
diet
#> # A tibble: 5,831 x 4
#>   binomial      diet_plant diet_vertibrate diet_invertebrate
#>   <chr>          <dbl>          <dbl>          <dbl>
#> 1 Abditomys_latidens      100              0              0
#> 2 Abeomelomys_sevia       78              3             19
#> 3 Abrawayaomys_ruschii     88              1             11
#> 4 Abrocoma_bennettii     100              0              0
#> 5 Abrocoma_boliviensis   100              0              0
#> 6 Abrocoma_budini        100              0              0
#> # ... with 5,825 more rows
```

651 These three variables show the percentage of each category of food that make
652 the diet of that species. They should sum to 100, unless the authors made a
653 typo or other entry error. To assert this, I'm going to create a new variable,
654 `total_diet`.

```
diet <- diet %>% mutate(
  "total_diet" = diet_vertibrate + diet_invertebrate + diet_plant
)
diet
#> # A tibble: 5,831 x 5
#>   binomial      diet_plant diet_vertibrate diet_invertebrate total_diet
#>   <chr>          <dbl>          <dbl>          <dbl>          <dbl>
#> 1 Abditomys_latidens      100              0              0            100
#> 2 Abeomelomys_sevia       78              3             19            100
#> 3 Abrawayaomys_ruschii     88              1             11            100
#> 4 Abrocoma_bennettii     100              0              0            100
#> 5 Abrocoma_boliviensis   100              0              0            100
```

```
#> 6 Abrocoma_budini          100          0          0          100
#> # ... with 5,825 more rows
```

```
all(diet$total_diet == 100)
#> [1] TRUE
# cool and good
```

655 `mutate()` adds a variable to the table, and keeps all other variables. Sometimes
 656 you may want to just keep the new variable, and drop the other ones. That's the
 657 job of `mutate()`'s twin sibling, `transmute()`. For example, I want to combine
 658 `diet_invertebrate` and `diet_vertibrate` together:

```
diet %>%
  transmute(
    "diet_animal" = diet_invertebrate + diet_vertibrate
  )
#> # A tibble: 5,831 x 1
#>   diet_animal
#>   <dbl>
#> 1          0
#> 2         22
#> 3         12
#> 4          0
#> 5          0
#> 6          0
#> # ... with 5,825 more rows
```

659 You may want to keep some variables and drop others. You could pipe
 660 `mutate()` and `select()` to do so, or you could just pass the variables to keep
 661 to `transmute()`.

```
diet %>%
  transmute(
    "diet_animal" = diet_invertebrate + diet_vertibrate,
    diet_plant
  )
#> # A tibble: 5,831 x 2
#>   diet_animal diet_plant
#>   <dbl>      <dbl>
#> 1          0        100
#> 2         22         78
#> 3         12         88
#> 4          0        100
#> 5          0        100
#> 6          0        100
#> # ... with 5,825 more rows
```

662 You can also refer to variables you're creating to derive new variables from them

663 as part of the same operation, this is not an issue.

```
diet %>%
  transmute(
    "diet_animal" = diet_invertebrate + diet_vertibrate,
    diet_plant,
    "total_diet" = diet_animal + diet_plant
  )
#> # A tibble: 5,831 x 3
#>   diet_animal diet_plant total_diet
#>   <dbl>      <dbl>      <dbl>
#> 1         0         100         100
#> 2        22         78         100
#> 3        12         88         100
#> 4         0        100         100
#> 5         0        100         100
#> 6         0        100         100
#> # ... with 5,825 more rows
```

664 Sometimes, you may need to perform an operation based on the row number (I
665 don't have a good example in mind). `tibble` has a built-in function to do just
666 that:

```
phylacine %>%
  select(binomial) %>%
  tibble::rownames_to_column(var = "row_nb")
#> # A tibble: 5,831 x 2
#>   row_nb binomial
#>   <chr> <chr>
#> 1 1      Abditomys_latidens
#> 2 2      Abeomelomys_sevia
#> 3 3      Abrawayaomys_ruschii
#> 4 4      Abrocoma_bennettii
#> 5 5      Abrocoma_boliviensis
#> 6 6      Abrocoma_budini
#> # ... with 5,825 more rows
```

667 3.4.2 Summarise observations with summarise()

668 `mutate()` applies operations to all observations in a table. By contrast,
669 `summarise()` applies operations to *groups* of observations, and returns, er,
670 summaries. The default grouping unit is the entire table:

```
phylacine %>%
  summarise(
    "nb_species" = n(), # counts observations
    "nb_terrestrial" = sum(terrestrial),
```

```

    "nb_marine" = sum(marine),
    "nb_freshwater" = sum(freshwater),
    "nb_aerial" = sum(aerial),
    "mean_mass_g" = mean(mass_g)
  )
#> # A tibble: 1 x 6
#>   nb_species nb_terrestrial nb_marine nb_freshwater nb_aerial mean_mass_g
#>   <int>          <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
#> 1      5831      4575       135       156       1162     156882.

```

671 Above you can see that bats account for a large portion of mammal species
 672 diversity (`nb_aerial`). How much exactly? Just as with `mutate()`, you can
 673 perform operations on the variables you just created, in the same statement:

```

phylacine %>%
  summarise(
    "nb_species" = n(),
    "nb_aerial" = sum(aerial), # bats
    "prop_aerial" = nb_aerial / nb_species
  )
#> # A tibble: 1 x 3
#>   nb_species nb_aerial prop_aerial
#>   <int>      <dbl>      <dbl>
#> 1      5831      1162      0.199

```

674 One fifth!

675 If the british spelling bothers you, `summarize()` exists and is strictly equivalent.

676 Here's a simple trick with logical (TRUE / FALSE) variables. Their sum is the
 677 count of observations that evaluate to TRUE (because TRUE is taken as 1 and
 678 FALSE as 0) and their mean is the proportion of TRUE observations. This can be
 679 exploited to count the number of observations that satisfy a condition:

```

phylacine %>%
  summarise(
    "nb_species" = n(),
    "nb_megafauna" = sum(mass_g > 100000),
    "p_megafauna" = mean(mass_g > 100000)
  )
#> # A tibble: 1 x 3
#>   nb_species nb_megafauna p_megafauna
#>   <int>      <int>      <dbl>
#> 1      5831       302      0.0518

```

680 There are more summaries that just means and counts (see `?summarise()` for
 681 some helpful functions). In fact, `summarise` can use any function or expression
 682 that evaluates to a single value or a *vector* of values. This includes base R `max()`,
 683 `quantiles`, etc.

684 `mutate()` and `transmute()` can compute summaries as well, but they will return
 685 the summary once for each observation, in a new column.

```

phylacine %>%
  mutate("nb_species" = n()) %>%
  select(binomial, nb_species)
#> # A tibble: 5,831 x 2
#>   binomial          nb_species
#>   <chr>              <int>
#> 1 Abditomys_latidens      5831
#> 2 Abeomelomys_sevia      5831
#> 3 Abrawayaomys_ruschii   5831
#> 4 Abrocoma_bennettii     5831
#> 5 Abrocoma_boliviensis   5831
#> 6 Abrocoma_budini        5831
#> # ... with 5,825 more rows

```

686 3.4.3 Grouping observations by variables

687 In most cases you don't want to run summary operations on the entire set of
 688 observations, but instead on observations that share a common value, i.e. groups.
 689 For example, I want to run the summary displayed above, but for each Order
 690 of mammals.

691 `distinct()` extracts all the unique values of a variable

```

phylacine %>% distinct(order)
#> # A tibble: 29 x 1
#>   order
#>   <chr>
#> 1 Rodentia
#> 2 Chiroptera
#> 3 Carnivora
#> 4 Pilosa
#> 5 Diprotodontia
#> 6 Cetartiodactyla
#> # ... with 23 more rows

```

692 I could work my way with what we have already seen, filtering observations
 693 (`filter(order == "Rodentia")`) and then pipeing the output to `summarise()`,
 694 and do it again for each Order. But that would be tedious.

695 Instead, I can use `group_by()` to pool observations by order.

```

phylacine %>%
  group_by(order)
#> # A tibble: 5,831 x 24
#> # Groups:   order [29]

```

```

#>   binomial order family genus species terrestrial marine freshwater aerial
#>   <chr>      <chr> <chr> <chr> <chr>          <dbl>  <dbl>          <dbl>  <dbl>
#> 1 Abditom~ Rode~ Murid~ Abdi~ latide~          1      0              0      0
#> 2 Abeomel~ Rode~ Murid~ Abeo~ sevia          1      0              0      0
#> 3 Abraway~ Rode~ Crice~ Abra~ ruschii        1      0              0      0
#> 4 Abrocom~ Rode~ Abroc~ Abro~ bennet~        1      0              0      0
#> 5 Abrocom~ Rode~ Abroc~ Abro~ bolivi~        1      0              0      0
#> 6 Abrocom~ Rode~ Abroc~ Abro~ budini         1      0              0      0
#> # ... with 5,825 more rows, and 15 more variables: life_habit_method <chr>,
#> #   life_habit_source <chr>, mass_g <dbl>, mass_method <chr>,
#> #   mass_source <chr>, mass_comparison <chr>, mass_comparison_source <chr>,
#> #   island_endemicity <chr>, iucn_status <chr>, added_iucn_status <chr>,
#> #   diet_plant <dbl>, diet_vertibrate <dbl>, diet_invertebrate <dbl>,
#> #   diet_method <chr>, diet_source <chr>

```

At first glance, nothing has changed, apart from an extra line of information in the output that tells me the observations have been grouped. But now here's what happen if I run the same `summarise()` statement on an ungrouped and a grouped table

```

phylacine %>%
  summarise(
    "n_species" = n(),
    "mean_mass_g" = mean(mass_g)
  )
#> # A tibble: 1 x 2
#>   n_species mean_mass_g
#>   <int>      <dbl>
#> 1     5831    156882.

phylacine %>%
  group_by(order) %>%
  summarise(
    "n_species" = n(),
    "mean_mass_g" = mean(mass_g)
  )
#> # A tibble: 29 x 3
#>   order          n_species mean_mass_g
#>   <chr>          <int>      <dbl>
#> 1 Afrosoricida         57        306.
#> 2 Carnivora           313     47905.
#> 3 Cetartiodactyla      392    1854811.
#> 4 Chiroptera          1162        49.1
#> 5 Cingulata            39     235529.
#> 6 Dasyuromorphia        74        748.
#> # ... with 23 more rows

```


700 I get one value for each group.

701 Observations can be grouped by multiple variables, which will output a summary
702 for every unique combination of groups.

```

phylacine %>%
  group_by(order, iucn_status) %>%
  summarise(
    "n_species" = n()
  )
#> # A tibble: 138 x 3
#> # Groups:   order [29]
#>   order      iucn_status n_species
#>   <chr>      <chr>      <int>
#> 1 Afrosoricida CR          1
#> 2 Afrosoricida DD          4
#> 3 Afrosoricida EN          7
#> 4 Afrosoricida EP          2
#> 5 Afrosoricida LC         32
#> 6 Afrosoricida NT          3
#> # ... with 132 more rows

```

703 Whenever you call `summarise()`, the last level of grouping is dropped. Note
704 how in the output table above, observations are still grouped by order, and no
705 longer by IUCN status. If I summarise observations again:

```

phylacine %>%
  group_by(order, iucn_status) %>%
  summarise(
    "n_species" = n()
  ) %>%
  summarise(
    "n_species_2" = n()
  )
#> # A tibble: 29 x 2
#>   order      n_species_2
#>   <chr>      <int>
#> 1 Afrosoricida          7
#> 2 Carnivora            8
#> 3 Cetartiodactyla       9
#> 4 Chiroptera            8
#> 5 Cingulata             5
#> 6 Dasyuromorphia        7
#> # ... with 23 more rows

```

706 I get the summary across orders, and the table is no longer grouped at all. This
707 is useful to consider if you need to work on summaries across different levels of
708 the data.

For example, I would like to know how the species in each order are distributed between the different levels of threat in the IUCN classification. To get these proportions, I need to first get the count of each number of species in a level of threat inside an order, and divide that by the number of species in that order.

```

phylacine %>%
  group_by(order, iucn_status) %>%
  summarise("n_order_iucn" = n()) %>%
  # grouping by iucn_status silently dropped
  mutate(
    "n_order" = sum(n_order_iucn),
    "p_iucn" = n_order_iucn / n_order
  )
#> # A tibble: 138 x 5
#> # Groups:   order [29]
#>   order      iucn_status n_order_iucn n_order p_iucn
#>   <chr>      <chr>          <int>   <int> <dbl>
#> 1 Afrosoricida CR              1      57 0.0175
#> 2 Afrosoricida DD              4      57 0.0702
#> 3 Afrosoricida EN              7      57 0.123
#> 4 Afrosoricida EP              2      57 0.0351
#> 5 Afrosoricida LC             32      57 0.561
#> 6 Afrosoricida NT              3      57 0.0526
#> # ... with 132 more rows

```

10.2% of Carnivores are Endangered (“EN”).

3.4.4 Grouped data and other dplyr verbs

Grouping does not only affect the behaviour of `summarise`, but under circumstances, other verbs can (and will!) perform operations by groups.

```

# Species with a higher mass than the mammal mean
phylacine %>%
  select("binomial", "mass_g") %>%
  filter(mass_g > mean(mass_g, na.rm = TRUE))
#> # A tibble: 234 x 2
#>   binomial      mass_g
#>   <chr>      <dbl>
#> 1 Alcelaphus_buselaphus 171002.
#> 2 Alces_alces          356998
#> 3 Archaeoindris_fontoynonti 160000
#> 4 Arctocephalus_pusillus 178500
#> 5 Arctodus_simus        709500
#> 6 Balaena_mysticetus    100000000
#> # ... with 228 more rows

```

```

# Species with a higher mass than the mean in their order
phylacine %>%
  group_by(order) %>%
  select("binomial", "mass_g") %>%
  filter(mass_g > mean(mass_g, na.rm = TRUE))
#> # A tibble: 890 x 3
#> # Groups:   order [27]
#>   order      binomial      mass_g
#>   <chr>      <chr>      <dbl>
#> 1 Chiroptera Acerodon_celebensis    390
#> 2 Chiroptera Acerodon_humilis      600.
#> 3 Chiroptera Acerodon_jubatus   1075
#> 4 Chiroptera Acerodon_leucotis    513.
#> 5 Chiroptera Acerodon_mackloti    470.
#> 6 Rodentia  Aeretes_melanopterus   732.
#> # ... with 884 more rows

# Largest mammal
phylacine %>%
  select(binomial, mass_g) %>%
  slice_max(mass_g)
#> # A tibble: 1 x 2
#>   binomial      mass_g
#>   <chr>      <dbl>
#> 1 Balaenoptera_musculus 190000000
# Largest species in each order
phylacine %>%
  group_by(order) %>%
  select(binomial, mass_g) %>%
  slice_max(mass_g)
#> # A tibble: 30 x 3
#> # Groups:   order [29]
#>   order      binomial      mass_g
#>   <chr>      <chr>      <dbl>
#> 1 Afrosoricida Plesiorycteropus_madagascariensis 13220
#> 2 Carnivora    Mirounga_leonina    1600000
#> 3 Cetartiodactyla Balaenoptera_musculus 190000000
#> 4 Chiroptera    Acerodon_jubatus    1075
#> 5 Cingulata     Glyptodon_clavipes  2000000
#> 6 Dasyuromorphia Thylacinus_cynocephalus 30000
#> # ... with 24 more rows

```

717 To avoid grouped operations, you can simply drop grouping with `ungroup()`.

718 3.5 Working with multiple tables

719 3.5.1 Binding tables

720 dplyr introduces `bind_rows()` and `bind_cols()`, which are equivalent to base
 721 R `rbind()` and `cbind()`, with a few extra feature. They are faster, and can
 722 bind many tables at once, and bind data frames with vectors or lists.

723 `bind_rows()` has an option to pass a variable specifying which dataset each
 724 observation originates from.

```
porpoises <- phylacine %>%
  filter(family == "Phocoenidae") %>%
  select(binomial, iucn_status)
echidnas <- phylacine %>%
  filter(family == "Tachyglossidae") %>%
  select(binomial, iucn_status)

bind_rows(
  "porpoise" = porpoises,
  "echidna" = echidnas,
  .id = "kind"
)
#> # A tibble: 13 x 3
#>   kind      binomial      iucn_status
#>   <chr>    <chr>          <chr>
#> 1 porpoise Neophocaena_asiaeorientalis VU
#> 2 porpoise Neophocaena_phocaenoides VU
#> 3 porpoise Phocoena_dioptrica      DD
#> 4 porpoise Phocoena_phocoena      LC
#> 5 porpoise Phocoena_sinus          CR
#> 6 porpoise Phocoena_spinipinnis    DD
#> # ... with 7 more rows
```

725 3.5.2 Combining variables of two tables with mutating 726 joins

727 Mutating joins are tailored to combine tables that share a set of observations
 728 but have different variables.

729 As an example, let's split the `phylacine` dataset in two smaller datasets, one
 730 containing information on diet and one on the dominant habitat.

```
diet <- phylacine %>%
  select(binomial, diet_plant:diet_invertebrate) %>%
  slice(1:5)
```

```
diet
#> # A tibble: 5 x 4
#>   binomial      diet_plant diet_vertbrate diet_invertebrate
#>   <chr>          <dbl>          <dbl>          <dbl>
#> 1 Abditomys_latidens      100            0            0
#> 2 Abeomelomys_sevia       78            3            19
#> 3 Abrawayaomys_ruschii     88            1            11
#> 4 Abrocoma_bennettii      100            0            0
#> 5 Abrocoma_boliviensis    100            0            0
```

```
life_habit <- phylacine %>% select(binomial, terrestrial:aerial) %>%
  slice(1:3, 6:7)
```

```
life_habit
#> # A tibble: 5 x 5
#>   binomial      terrestrial marine freshwater aerial
#>   <chr>          <dbl>   <dbl>      <dbl>   <dbl>
#> 1 Abditomys_latidens      1     0        0     0
#> 2 Abeomelomys_sevia      1     0        0     0
#> 3 Abrawayaomys_ruschii    1     0        0     0
#> 4 Abrocoma_budini        1     0        0     0
#> 5 Abrocoma_cinerea       1     0        0     0
```

731 The two datasets each contain 5 species, the first three are shared, and the two
732 last differ between the two.

```
intersect(diet$binomial, life_habit$binomial)
#> [1] "Abitomys_latidens" "Abeomelomys_sevia" "Abrawayaomys_ruschii"
setdiff(diet$binomial, life_habit$binomial)
#> [1] "Abrocoma_bennettii" "Abrocoma_boliviensis"
```

733 To use mutate-joins, both tables need to have a **key**, a variable that identifies
734 each observation. Here, that would be `binomial`, the species names. If
735 your table doesn't have such a key and the rows between the tables match
736 one another, remember you can create a row number variable easily with
737 `tibble::column_to_rownames()`.

```
inner_join(diet, life_habit, by = "binomial")
#> # A tibble: 3 x 8
#>   binomial diet_plant diet_vertbrate diet_invertebra~ terrestrial marine
#>   <chr>      <dbl>          <dbl>          <dbl>      <dbl>   <dbl>
#> 1 Abditom~      100            0            0          1     0
#> 2 Abeomel~       78            3            19          1     0
#> 3 Abraway~       88            1            11          1     0
#> # ... with 2 more variables: freshwater <dbl>, aerial <dbl>
```

738 `inner_join` combined the variables, and dropped the observations that weren't
739 matched between the two tables. There are three other variations of mutating
740 joins, differing in what they do with unmatched variables.

```

left_join(diet, life_habit, by = "binomial")
#> # A tibble: 5 x 8
#>   binomial diet_plant diet_vertibrate diet_invertebra~ terrestrial marine
#>   <chr>      <dbl>          <dbl>          <dbl>          <dbl> <dbl>
#> 1 Abditom~    100            0            0            1     0
#> 2 Abeomel~    78            3            19           1     0
#> 3 Abraway~    88            1            11           1     0
#> 4 Abrocom~   100            0            0           NA     NA
#> 5 Abrocom~   100            0            0           NA     NA
#> # ... with 2 more variables: freshwater <dbl>, aerial <dbl>
right_join(diet, life_habit, by = "binomial")
#> # A tibble: 5 x 8
#>   binomial diet_plant diet_vertibrate diet_invertebra~ terrestrial marine
#>   <chr>      <dbl>          <dbl>          <dbl>          <dbl> <dbl>
#> 1 Abditom~    100            0            0            1     0
#> 2 Abeomel~    78            3            19           1     0
#> 3 Abraway~    88            1            11           1     0
#> 4 Abrocom~    NA            NA            NA           1     0
#> 5 Abrocom~    NA            NA            NA           1     0
#> # ... with 2 more variables: freshwater <dbl>, aerial <dbl>
full_join(diet, life_habit, by = "binomial")
#> # A tibble: 7 x 8
#>   binomial diet_plant diet_vertibrate diet_invertebra~ terrestrial marine
#>   <chr>      <dbl>          <dbl>          <dbl>          <dbl> <dbl>
#> 1 Abditom~    100            0            0            1     0
#> 2 Abeomel~    78            3            19           1     0
#> 3 Abraway~    88            1            11           1     0
#> 4 Abrocom~    100            0            0           NA     NA
#> 5 Abrocom~    100            0            0           NA     NA
#> 6 Abrocom~    NA            NA            NA           1     0
#> # ... with 1 more row, and 2 more variables: freshwater <dbl>, aerial <dbl>

semi_join(diet, life_habit, by = "binomial")
#> # A tibble: 3 x 4
#>   binomial diet_plant diet_vertibrate diet_invertebrate
#>   <chr>      <dbl>          <dbl>          <dbl>
#> 1 Abditomys_latidens    100            0            0
#> 2 Abeomelomys_sevia     78            3            19
#> 3 Abrawayaomys_ruschii  88            1            11
anti_join(diet, life_habit, by = "binomial")
#> # A tibble: 2 x 4
#>   binomial diet_plant diet_vertibrate diet_invertebrate
#>   <chr>      <dbl>          <dbl>          <dbl>
#> 1 Abrocoma_bennettii    100            0            0
#> 2 Abrocoma_bolivienis  100            0            0

```

741 3.5.3 Filtering matching observations between two tables 742 wiht filtering joins

743 So-called filtering joins return row from the first table that are matched (or not,
744 for `anti_join()`) in the second.

```

semi_join(diet, life_habit, by = "binomial")
#> # A tibble: 3 x 4
#>   binomial      diet_plant diet_vertibrate diet_invertebrate
#>   <chr>          <dbl>          <dbl>          <dbl>
#> 1 Abditomys_latidens      100            0            0
#> 2 Abeomelomys_sevia       78            3           19
#> 3 Abrawayaomys_ruschii     88            1           11
anti_join(diet, life_habit, by = "binomial")
#> # A tibble: 2 x 4
#>   binomial      diet_plant diet_vertibrate diet_invertebrate
#>   <chr>          <dbl>          <dbl>          <dbl>
#> 1 Abrocoma_bennettii      100            0            0
#> 2 Abrocoma_boliviensis    100            0            0

```


Chapter 4

Working with lists and iteration

Every use case is ridiculous
until it happens to you.

```
# load the tidyverse  
library(tidyverse)
```

4.1 List columns with tidyr

4.1.1 Nesting data

It may become necessary to indicate the groups of a tibble in a somewhat more explicit way than simply using `dplyr::group_by`. `tidyr` offers the option to

753 create nested tibbles, that is, to store complex objects in the columns of a tibble.
 754 This includes other tibbles, as well as model objects and plots.

755 *NB:* Nesting data is done using `tidyr::nest`, which is different from the simi-
 756 larly named `tidyr::nesting`.

757 The example below shows how `mtcars` can be converted into a nested tibble.

```
# nest mtcars into a list of dataframes based on number of cylinders
nested_cars = as_tibble(mtcars,
                        rownames = "car_name") %>%
  group_by(cyl) %>%
  nest()

nested_cars
#> # A tibble: 3 x 2
#> # Groups:   cyl [3]
#>   cyl data
#>   <dbl> <list>
#> 1     6 <tibble [7 x 11]>
#> 2     4 <tibble [11 x 11]>
#> 3     8 <tibble [14 x 11]>

# get column class
sapply(nested_cars, class)
#>      cyl      data
#> "numeric"  "list"
```

758 `mtcars` is now a nested data frame. The class of each of its columns is respec-
 759 tively, a numeric (number of cylinders) and a list (the data of all cars with as
 760 many cylinders as in the corresponding row).

761 While `nest` can be used without first grouping the tibble, it's just much easier
 762 to group first.

763 4.1.2 Unnesting data

764 A nested tibble can be converted back into the original, or into a processed form,
 765 using `tidyr::unnest`. The original groups are retained.

```
# use unnest to recover the original data frame
unnest(nested_cars, cols = "data")
#> # A tibble: 32 x 12
#> # Groups:   cyl [3]
#>   cyl car_name      mpg disp   hp  drat   wt  qsec   vs   am gear carb
#>   <dbl> <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1     6 Mazda RX4      21  160  110  3.9  2.62  16.5    0    1    4    4
#> 2     6 Mazda RX4 W~  21  160  110  3.9  2.88  17.0    0    1    4    4
```

```
#> 3      6 Hornet 4 Dr~ 21.4 258    110 3.08 3.22 19.4    1    0    3    1
#> 4      6 Valiant      18.1 225    105 2.76 3.46 20.2    1    0    3    1
#> 5      6 Merc 280      19.2 168.   123 3.92 3.44 18.3    1    0    4    4
#> 6      6 Merc 280C     17.8 168.   123 3.92 3.44 18.9    1    0    4    4
#> # ... with 26 more rows

# unnesting preserves groups
groups(unnest(nested_cars, cols = "data"))
#> [[1]]
#> cyl
```

766 The `unnest_longer` and `unnest_wider` variants of `unnest` are maturing func-
 767 tions, that is, not in their final form. They allow interesting variations on
 768 unnesting — these are shown here but advised against.

769 Unnest the data first, and then convert it to the form needed.

```
unnest_longer(nested_cars, col = "data") %>%
  head()
#> # A tibble: 6 x 2
#> # Groups:   cyl [1]
#>   cyl data$car_name $mpg $disp $hp $drat $wt $qsec $vs $am $gear
#>   <dbl> <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1      6 Mazda RX4      21 160    110 3.9 2.62 16.5    0    1    4
#> 2      6 Mazda RX4 Wag  21 160    110 3.9 2.88 17.0    0    1    4
#> 3      6 Hornet 4 Dri~ 21.4 258    110 3.08 3.22 19.4    1    0    3
#> 4      6 Valiant      18.1 225    105 2.76 3.46 20.2    1    0    3
#> 5      6 Merc 280      19.2 168.   123 3.92 3.44 18.3    1    0    4
#> 6      6 Merc 280C     17.8 168.   123 3.92 3.44 18.9    1    0    4
#> # ... with 1 more variable: $carb <dbl>

unnest_wider(nested_cars, col = "data")
#> # A tibble: 3 x 12
#> # Groups:   cyl [3]
#>   cyl car_name mpg disp hp drat wt qsec vs am gear carb
#>   <dbl> <list> <list> <list> <list> <lis> <lis> <lis> <lis> <lis> <lis> <lis>
#> 1      6 <chr [7]> <dbl ~ <dbl ~ <dbl ~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~
#> 2      4 <chr [11~ <dbl ~ <dbl ~ <dbl ~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~
#> 3      8 <chr [14~ <dbl ~ <dbl ~ <dbl ~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~ <dbl~
```

770 4.1.3 Working with list columns

771 The class of a list column is `list`, and working with list columns (and lists, and
 772 list-like objects such as vectors) makes iteration necessary, since this is one of
 773 the only ways to operate on lists.

Two examples are shown below when getting the class and number of rows of the nested tibbles in the list column.

```
# how many rows in each nested tibble?
for (i in seq_along(nested_cars$data)) {
  print(nrow(nested_cars$data[[i]]))
}
#> [1] 7
#> [1] 11
#> [1] 14

# what is the class of each element?
lapply(X = nested_cars$data, FUN = class)
#> [[1]]
#> [1] "tbl_df"      "tbl"        "data.frame"
#>
#> [[2]]
#> [1] "tbl_df"      "tbl"        "data.frame"
#>
#> [[3]]
#> [1] "tbl_df"      "tbl"        "data.frame"
```

Functionals

The second example uses `lapply`, and this is a *functional*. *Functionals* are functions that take another function as one of their arguments. Base R functionals include the `*apply` family of functions: `apply`, `lapply`, `vapply` and so on.

4.2 Iteration with map

The `tidyverse` replaces traditional loop-based iteration with *functionals* from the `purrr` package. A good reason to use `purrr` functionals instead of base R functionals is their consistent and clear naming, which always indicates how they should be used. This is explained in the examples below.

How `map` is different from `for` and `lapply` are best explained in the **Advanced R Book**.

4.2.1 Basic use of map

`map` works very similarly to `lapply`, where `.x` is object on whose elements to apply the function `.f`.

```
# get the number of rows in data
map(.x = nested_cars$data, .f = nrow)
#> [[1]]
#> [1] 7
#>
#> [[2]]
#> [1] 11
#>
#> [[3]]
#> [1] 14
```

790 `map` works on any list-like object, which includes vectors, and always returns a
 791 list. `map` takes two arguments, the object on which to operate, and the function
 792 to apply to each element.

```
# get the square root of each integer 1 - 10
some_numbers = 1:3
map(some_numbers, sqrt)
#> [[1]]
#> [1] 1
#>
#> [[2]]
#> [1] 1.41
#>
#> [[3]]
#> [1] 1.73
```

793 4.2.2 map variants returning vectors

794 Though `map` always returns a list, it has variants named `map_*` where the suffix
 795 indicates the return type. `map_chr`, `map_dbl`, `map_int`, and `map_lgl` return
 796 character, double (numeric), integer, and logical vectors.

```
# use map_dbl to get a vector of square roots
some_numbers = 1:10
map_dbl(some_numbers, sqrt)
#> [1] 1.00 1.41 1.73 2.00 2.24 2.45 2.65 2.83 3.00 3.16

# map_chr will convert the output to a character
map_chr(some_numbers, sqrt)
#> [1] "1.000000" "1.414214" "1.732051" "2.000000" "2.236068" "2.449490"
#> [7] "2.645751" "2.828427" "3.000000" "3.162278"

# map_lgl returns TRUE/FALSE values
some_numbers = c(NA, 1:3, NA, NaN, Inf, -Inf)
map_lgl(some_numbers, is.na)
```

```
#> [1] TRUE FALSE FALSE FALSE TRUE TRUE FALSE FALSE
```

797 4.2.3 map variants returning data frames

798 `map_df` returns data frames, and by default binds dataframes by rows, while
799 `map_dfr` does this explicitly, and `map_dfc` does returns a dataframe bound by
800 column.

```
# split mtcars into 3 dataframes, one per cylinder number
some_list = split(mtcars, mtcars$cyl)

# get the first two rows of each dataframe
map_df(some_list, head, n = 2)
#>           mpg cyl disp  hp drat   wt  qsec vs am gear carb
#> Datsun 710   22.8   4  108  93 3.85 2.32 18.6  1  1    4    1
#> Merc 240D   24.4   4  147  62 3.69 3.19 20.0  1  0    4    2
#> Mazda RX4   21.0   6  160 110 3.90 2.62 16.5  0  1    4    4
#> Mazda RX4 Wag 21.0   6  160 110 3.90 2.88 17.0  0  1    4    4
#> Hornet Sportabout 18.7   8  360 175 3.15 3.44 17.0  0  0    3    2
#> Duster 360   14.3   8  360 245 3.21 3.57 15.8  0  0    3    4
```

801 `map` accepts arguments to the function being mapped, such as in the example
802 above, where `head()` accepts the argument `n = 2`.

803 `map_dfr` behaves the same as `map_df`.

```
# the same as above but with a pipe
some_list %>%
  map_dfr(head, n = 2)
#>           mpg cyl disp  hp drat   wt  qsec vs am gear carb
#> Datsun 710   22.8   4  108  93 3.85 2.32 18.6  1  1    4    1
#> Merc 240D   24.4   4  147  62 3.69 3.19 20.0  1  0    4    2
#> Mazda RX4   21.0   6  160 110 3.90 2.62 16.5  0  1    4    4
#> Mazda RX4 Wag 21.0   6  160 110 3.90 2.88 17.0  0  1    4    4
#> Hornet Sportabout 18.7   8  360 175 3.15 3.44 17.0  0  0    3    2
#> Duster 360   14.3   8  360 245 3.21 3.57 15.8  0  0    3    4
```

804 `map_dfc` binds the resulting 3 data frames of two rows each by column, and
805 automatically repairs the column names, adding a suffix to each duplicate.

```
some_list %>%
  map_dfc(head, n = 2)
#>           mpg...1 cyl...2 disp...3 hp...4 drat...5 wt...6 qsec...7 vs...8
#> Datsun 710   22.8     4    108    93   3.85   2.32   18.6     1
#> Merc 240D   24.4     4    147    62   3.69   3.19   20.0     1
#>           am...9 gear...10 carb...11 mpg...12 cyl...13 disp...14 hp...15
#> Datsun 710     1     4          1    21     6    160    110
#> Merc 240D     0     4          2    21     6    160    110
```

```

#>          drat...16 wt...17 qsec...18 vs...19 am...20 gear...21 carb...22
#> Datsun 710      3.9   2.62    16.5      0      1      4      4
#> Merc 240D      3.9   2.88    17.0      0      1      4      4
#>          mpg...23 cyl...24 disp...25 hp...26 drat...27 wt...28 qsec...29
#> Datsun 710    18.7      8    360    175    3.15    3.44    17.0
#> Merc 240D    14.3      8    360    245    3.21    3.57    15.8
#>          vs...30 am...31 gear...32 carb...33
#> Datsun 710      0      0      3      2
#> Merc 240D      0      0      3      4

```

806 4.2.4 Working with list columns using map

807 The various `map` versions integrate well with list columns to make syn-
 808 thetic/summary data. In the example, the `dplyr::mutate` function is used to
 809 add three columns to the nested tibble: the number of rows, the mean mileage,
 810 and the name of the first car.

811 In each of these cases, the vectors added are generated using `purrr` functions.

```

# get the number of rows per dataframe, the mean mileage, and the first car
nested_cars = nested_cars %>%
  mutate(
    # use the int return to get the number of rows
    n_rows = map_int(data, nrow),

    # double return for mean mileage
    mean_mpg = map_dbl(data, function(df) {mean(df$mpg)}),

    # character return to get first car
    first_car = map_chr(data, function(df) {first(df$car_name)})
  )

# examine the output
nested_cars
#> # A tibble: 3 x 5
#> # Groups:   cyl [3]
#>   cyl data          n_rows mean_mpg first_car
#>   <dbl> <list>      <int>   <dbl> <chr>
#> 1     6 <tibble [7 x 11]>      7    19.7 Mazda RX4
#> 2     4 <tibble [11 x 11]>     11    26.7 Datsun 710
#> 3     8 <tibble [14 x 11]>     14    15.1 Hornet Sportabout

```

812 4.2.5 Selective mapping using map variants

813 `map_at` and `map_if` work like other `*_at` and `*_if` functions. Here, `map_if` is
 814 used to run a linear model only on those tibbles which have sufficient data. The
 815 predicate is specified by `.p`.

816 In this example, the nested tibble is given a new column using `dplyr::mutate`,
 817 where the data to be added is a mixed list.

```
# split mtcars by cylinder number and run an lm only if there are more than 10 rows
data = nest(mtcars, data = -cyl)

data = mutate(data,
               model = map_if(.x = data,
                             .p = function(x){
                               nrow(x) > 10
                             },
                             .f = function(x){
                               lm(mpg ~ wt, data = x)
                             })))

# check the data structure
data
#> # A tibble: 3 x 3
#>   cyl data          model
#>   <dbl> <list>         <list>
#> 1     6 <tibble [7 x 10]> <tibble [7 x 10]>
#> 2     4 <tibble [11 x 10]> <lm>
#> 3     8 <tibble [14 x 10]> <lm>
```

818 The first element is a tibble of the corresponding element in `mtcars$cars`, which
 819 has not been operated on because it has fewer than 10 rows. The remaining
 820 elements are `lm` objects.

821 4.3 More map variants

822 `map` also has variants along the axis of how many elements are operated upon.
 823 `map2` operates on two vectors or list-like elements, and returns a single list as
 824 output, while `pmap` operates on a list of list-like elements. The output has as
 825 many elements as the input lists, which must be of the same length.

826 4.3.1 Mapping over two inputs with map2

827 `map2` has the same variants as `map`, allowing for different return types. Here
 828 `map2_int` returns an integer vector.


```
# consider 2 vectors and replicate the simple vector addition using map2
map2_int(.x = 1:5,
        .y = 6:10,
        .f = sum)
#> [1]  7  9 11 13 15
```

829 `map2` doesn't have `_at` and `_if` variants.

830 One use case for `map2` is to deal with both a list element and its index, as shown
 831 in the example. This may be necessary when the list index is removed in a
 832 `split` or `nest`. This can also be done with `imap`, where the index is referred to
 833 as `.y`.

```
# make a named list for this example
this_list = list(a = "first letter",
                b = "second letter")
```

```
# a not particularly useful example
map2(this_list, names(this_list),
     function(x, y) {
       glue::glue('{x} : {y}')
     })
```

```
#> $a
#> first letter : a
#>
#> $b
#> second letter : b
```

```
# imap can also do this
imap(this_list,
     function(x, .y){
       glue::glue('{x} : {.y}')
```

```
#> $a
#> first letter : a
#>
#> $b
#> second letter : b
```

834 4.3.2 Mapping over multiple inputs with `pmap`

835 `pmap` instead operates on a list of multiple list-like objects, and also comes with
 836 the same return type variants as `map`. The example shows both aspects of `pmap`
 837 using `pmap_chr`.

```
# operate on three different lists
list_01 = as.list(1:3)
```

```

list_02 = as.list(letters[1:3])
list_03 = as.list(rainbow(3))

# print a few statements
pmap_chr(list(list_01, list_02, list_03),
  function(l1, l2, l3){
    glue::glue('number {l1}, letter {l2}, colour {l3}')
  })
#> [1] "number 1, letter a, colour #FF0000" "number 2, letter b, colour #00FF00"
#> [3] "number 3, letter c, colour #0000FF"

```

838 4.3.3 Mapping at depth

839 Lists are often nested, that is, a list element may itself be a list. It is possible
 840 to map a function over elements as a specific depth.

841 In the example, `mtcars` is split by cylinders, and then by gears, creating a
 842 two-level list, with the second layer operated on.

```

# use map to make a 2 level list
this_list = split(mtcars, mtcars$cyl) %>%
  map(function(df){ split(df, df$gear) })

# map over the second level to count the number of
# cars with N gears in the set of cars with M cylinders
# display only for cyl = 4
map_depth(this_list[1], 2, nrow)
#> $`4`
#> $`4`$`3`
#> [1] 1
#>
#> $`4`$`4`
#> [1] 8
#>
#> $`4`$`5`
#> [1] 2

```

843 4.3.4 Iteration without a return

844 `map` and its variants have a return type, which is either a list or a vector. How-
 845 ever, it is often necessary to iterate a function over a list-like object for that
 846 function's side effects, such as printing a message to screen, plotting a series of
 847 figures, or saving to file.

848 `walk` is the function for this task. It has only the variants `walk2`, `iwalk`, and

849 `pwalk`, whose logic is similar to `map2`, `imap`, and `pmap`. In the example, the
 850 function applied to each list element is intended to print a message.

```
this_list = split(mtcars, mtcars$cyl)

iwalk(this_list,
  function(df, .y){
    message(glue::glue('{nrow(df)} cars with {.y} cylinders'))
  })
```

851 4.3.5 Modify rather than map

852 When the return type is expected to be the same as the input type, that is, a
 853 list returning a list, or a character vector returning the same, `modify` can help
 854 with keeping strictly to those expectations.

855 In the example, simply adding 2 to each vector element produces an error,
 856 because the output is a `numeric`, or `double`. `modify` helps ensure some type
 857 safety in this way.

```
vec = as.integer(1:10)

tryCatch(
  expr = {

    # this is what we want you to look at

    modify(vec, function(x) { (x + 2) })

  },

  # do not pay attention to this
  error = function(e){
    print(toString(e))
  }
)
#> [1] "Error: Can't coerce element 1 from a double to a integer\n"
```

858 Converting the output to an integer, which was the original input type, serves
 859 as a solution.

```
modify(vec, function(x) { as.integer(x + 2) })
#> [1] 3 4 5 6 7 8 9 10 11 12
```

860 A note on `invoke`

861 `invoke` used to be a wrapper around `do.call`, and can still be found with its
 862 family of functions in `purrr`. It is however retired in favour of functionality
 863 already present in `map` and `rlang::exec`, the latter of which will be covered in
 864 another session.

865 4.4 Other functions for working with lists

866 `purrr` has a number of functions to work with lists, especially lists that are not
 867 nested list-columns in a tibble.

868 4.4.1 Filtering lists

869 Lists can be filtered on any predicate using `keep`, while the special case `compact`
 870 is applied when the empty elements of a list are to be filtered out. `discard` is
 871 the opposite of `keep`, and keeps only elements not satisfying a condition. Again,
 872 the predicate is specified by `.p`.

```
# a list containing numbers
this_list = list(a = 1, b = -1, c = 2, d = NULL, e = NA)

# remove the empty element
# this must be done before using keep on the list
this_list = compact(this_list)

# use discard to remove the NA
this_list = discard(this_list, .p = is.na)

# keep list elements which are positive
keep(this_list, .p = function(x){ x > 0 })
#> $a
#> [1] 1
#>
#> $c
#> [1] 2
```

873 `head_while` is bit of an odd case, which returns all elements of a list-like object
 874 in sequence until the first one fails to satisfy a predicate, specified by `.p`.

```
1:10 %>%
  head_while(.p = function(x) x < 5)
#> [1] 1 2 3 4
```

875 4.4.2 Summarising lists

876 The `purrr` functions `every`, `some`, `has_element`, `detect`, `detect_index`, and
 877 `vec_depth` help determine whether a list passes a certain logical test or not.
 878 These are seldom used and are not discussed here.

879 4.4.3 Reduction and accumulation

880 `reduce` helps combine elements along a list using a specific function. Consider
 881 the example below where list elements are concatenated into a single vector.

```
this_list = list(a = 1:3, b = 3:4, c = 5:10)
```

```
reduce(this_list, c)
#> [1] 1 2 3 3 4 5 6 7 8 9 10
```

882 This can also be applied to data frames. Consider some random samples of
 883 `mtcars`, each with only 5 cars removed. The objective is to find the cars present
 884 in all 10 samples.

885 The way `reduce` works in the example below is to take the first element and find
 886 its intersection with the second, and to take the result and find its intersection
 887 with the third and so on.

```
# sample mtcars
mtcars = as_tibble(mtcars, rownames = "car")
sampled_data = map(1:10, function(x){sample_n(mtcars, nrow(mtcars)-5)})

# get cars which appear in all samples
sampled_data = reduce(sampled_data, dplyr::inner_join)
```

888 `accumulate` works very similarly, except it retains the intermediate products.
 889 The first element is retained as is. `accumulate2` and `reduce2` work on two lists,
 890 following the same logic as `map2` etc. Both functions can be used in much more
 891 complex ways than demonstrated here.

```
# make a list
this_list = list(a = 1:3, b = 3:6, c = 5:10, d = c(1,2,5,10,12))

# a multiple accumulate can help
accumulate(this_list, union, .dir = "forward")
#> $a
#> [1] 1 2 3
#>
#> $b
#> [1] 1 2 3 4 5 6
#>
#> $c
```

```
#> [1] 1 2 3 4 5 6 7 8 9 10
#>
#> $d
#> [1] 1 2 3 4 5 6 7 8 9 10 12
```

892 4.4.4 Miscellaneous operation

893 **purrr** offers a few more functions to work with lists (or list like objects).
 894 **prepend** works very similarly to **append**, except it adds to the head of a list.
 895 **splice** adds multiple objects together in a list. **splice** will break the existing
 896 list structure of input lists.

```
# use prepend to add values to the head of a list
prepend(x = list("a", "b"), values = list("1", "2"))
#> [[1]]
#> [1] "1"
#>
#> [[2]]
#> [1] "2"
#>
#> [[3]]
#> [1] "a"
#>
#> [[4]]
#> [1] "b"

# use splice to add multiple elements together
splice(list("a", "b"), list("1", "2"), "something else")
#> [[1]]
#> [1] "a"
#>
#> [[2]]
#> [1] "b"
#>
#> [[3]]
#> [1] "1"
#>
#> [[4]]
#> [1] "2"
#>
#> [[5]]
#> [1] "something else"
```

897 **flatten** has a similar behaviour, and converts a list of vectors or list of lists to a
 898 single list-like object. **flatten_*** options allow the output type to be specified.

```
this_list = list(a = rep("a", 3),
                 b = rep("b", 4))
```

```
this_list
#> $a
#> [1] "a" "a" "a"
#>
#> $b
#> [1] "b" "b" "b" "b"
```

```
# use flatten_chr to get a character vector
flatten_chr(this_list)
#> [1] "a" "a" "a" "b" "b" "b" "b"
```

899 **transpose** shifts the index order in multi-level lists. This is seen in the example,
 900 where the **gear** goes from being the index of the second level to the index of the
 901 first.

```
this_list = split(mtcars, mtcars$cyl) %>%
  map(function(df) split(df, df$gear))
```

```
# from a list of lists where cars are divided by cylinders and then
# gears, this is now a list of lists where cars are divided by
# gears and then cylinders
transpose(this_list[1])
#> $`3`
#> $`3`$`4`
#> # A tibble: 1 x 12
#>   car      mpg   cyl  disp    hp  drat    wt   qsec    vs    am  gear  carb
#>   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1 Toyota Coro~ 21.5     4  120.   97   3.7  2.46  20.0     1     0     3     1
#>
#>
#> $`4`
#> $`4`$`4`
#> # A tibble: 8 x 12
#>   car      mpg   cyl  disp    hp  drat    wt   qsec    vs    am  gear  carb
#>   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1 Datsun 710   22.8     4  108    93   3.85  2.32  18.6     1     1     4     1
#> 2 Merc 240D   24.4     4  147    62   3.69  3.19  20.0     1     0     4     2
#> 3 Merc 230    22.8     4  141    95   3.92  3.15  22.9     1     0     4     2
#> 4 Fiat 128    32.4     4   78.7    66   4.08  2.2   19.5     1     1     4     1
#> 5 Honda Civic 30.4     4   75.7    52   4.93  1.62  18.5     1     1     4     2
#> 6 Toyota Coro~ 33.9     4   71.1    65   4.22  1.84  19.9     1     1     4     1
#> # ... with 2 more rows
#>
#>
```

```

#> $`5`
#> $`5`$`4`
#> # A tibble: 2 x 12
#>   car      mpg   cyl  disp    hp  drat    wt   qsec    vs    am  gear  carb
#>   <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1 Porsche 914~ 26     4 120.    91  4.43  2.14  16.7    0     1     5     2
#> 2 Lotus Europa 30.4    4  95.1   113  3.77  1.51  16.9    1     1     5     2

```

902 4.5 To add: patchwork

903 4.5.0.1 Final words

904 In general, an iteration based problem can usually be solved with `purrr`.